

# 資料處理方法

平滑技巧/遺失值處理  
資料轉換/重抽法則

吳漢銘

國立臺北大學 統計學系

## ■ 主題1 [進階選讀]

- 移動平均 (Moving Average)
- 曲線配適 (Fitting Curves): lowess
- 核密度估計 (Kernel Density Estimation)
- 三次樣條插值 (Cubic Spline Interpolation)

## ■ 主題2

- 具缺失值資料 (Missing Data)
- 缺失機制 (Missingness Mechanism)
  - Missing by Design
  - Missing Completely at Random (MCAR)
  - Missing at Random (MAR)
  - Missing Not at Random (MNAR)

## ■ 主題3

- R Packages for Dealing With Missing Values: VIM, MICE
- Visualizing the Pattern of Missing Data
- Traditional Approaches to Handling Missing Data
- Imputation Methods: KNN
- Which Imputation Method?

## ■ 主題4

- 為什麼要做資料轉換?
- 常見的資料轉換方式
- 對數轉換 (Log Transformation)
- Box-Cox Transformation
- 標準化 (Standardization)
- 要使用哪一種資料轉換方式?

## ■ 主題5

- Training data and Testing data
- Resampling methods
  - Jackknife (leave-one-out)
  - Bootstrapping
- Ensemble Learning
  - bagging
  - boosting

## ■ 主題6

- Imbalanced Data Problem
  - under-sampling
  - over-sampling

# 具缺失值資料 (Missing Data)

3/70

Missing data (missing values for **certain variables for certain cases**): **item non-response**.

When data are missing for **a variable for all cases**: **latent or unobserved**.

When data are missing for **all variables for a given case**: **unit non-response**.

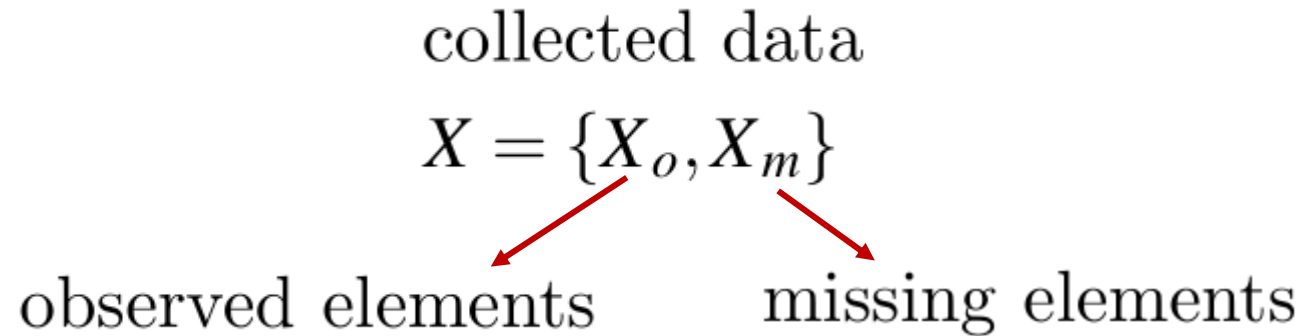
	A	B	C	D	E	F	G
1	ID	C	Y	X1	X2	X3	X4
2	s1	1	78.3	69.6	74.3	NA	5.22
3	s2	2	77	69.9	72.54	NA	3.98
4	s3	3	72.2	65.7	69.74	NA	4.89
5	s4	1	33.4	NA	30.97	NA	21.54
6	s5	2	32.65	28.35	30.54	NA	9.82
7	s6	3	35.45	28.5	32.01	NA	19.81
8	s7	1	424	378	403.55	NA	12.98
9	s8	2	NA	NA	NA	NA	NA
10	s9	3	355	312.5	339.96	NA	14.14
11	s10	1	18.2	15.5	17.19	NA	13.93
12	s11	2	18.3	15.3	16.38	NA	6.92
13	s12	3	16.1	13.9	14.92	NA	10.15
14	s13	1	23.75	20.2	22.19	NA	32.81

- The missing values may give clues to **systematic aspects of the problem.**
- How to deal with missing values:
  - Use a **global constant** to fill the value will misguide the mining process. (例如: 缺考給0分; 影像訊號=前景-背景)
  - Use the **attribute mean** or **median** for all samples belonging to the **same class** as the given tuple.
  - 補值 (Missing value imputation) (most popular)

## Missingness Mechanism

- The presence of missing data can
  - effect the properties of the estimates  
(e.g. means, percentages, percentiles, variances, ratios, regression parameters, etc.).
  - affect inferences.  
(e.g., the properties of tests and confidence intervals.)
- The missingness mechanism (Little and Rubin, 1987)
  - The way in which the probability of an item missing depends on other observed or non-observed variables as well as on its own value.
- It helpful to classify missing values on the basis of the stochastic mechanism that produces them.

# Missingness Mechanism



The missingness indicator matrix  $R$  corresponds  $X$ ,  
and each element of  $R$  is 1 if the corresponding element of  $X$  is missing,  
and 0 otherwise.

define the missingness mechanism as  
the probability of  $R$  conditional on  
the values of the observed and missing elements of  $X$ :

$$Pr(R|X_o, X_m)$$



# Missing by Design

## Missing Completely at Random

7/70

- Missing by Design

- Excluded some participants from the analysis because they are **not part of** the population under investigation.
- **missingness codes:** (i) refused to answer; (ii) answered don't know; (iii) had a valid skip or (iv) was skipped by an enumerator error.

- Missing Completely at Random (MCAR)

- missingness is **independent** of their own unobserved values and the observed data.

$$Pr(R|X) = Pr(R)$$

- **Example:** Miscoding or forgetting to log in answer.
- Imputation methods rely on the missingness being of the **MCAR** type.



# Missing at Random (MAR) Missing Not at Random (MNAR)

8/70

- Missing at Random (MAR)  $Pr(R|X) = Pr(R|X_o)$ 
  - missingness does not depend on their **unobserved value** but does depend on the **observed data**.
  - **Example 1**: male participants (**observed data**) are more likely to refuse to fill out the **depression survey**, but it does not depend on the level of their depression (**unobserved value**).
  - **Example 2**: if men are more likely to tell you their weight than women, **weight** is MAR.
  - We **can ignore missing data** (= omit missing observations) if we have **MAR** or **MCAR**.
- Missing Not at Random (MNAR)
  - Missingness that **depends on the missing** value itself.
  - **Example**: question about **income**, where the high rate of missing values (usually 20%~50%) is related to the value of the income itself (**very high and very low values will not be answered**).
  - **MNAR data is a more serious issue. (not ignorable)**



# Some Notes

- Assuming data is **MCAR**, too much missing data can be a problem.
  - Usually a safe maximum threshold is **5%** of the total for large datasets.
  - If missing data for a certain feature or sample is more than **5%** then you probably should leave that feature or sample **out**.
- If some variable is missing almost **25%** of the data points.
  - Consider either dropping it from the analysis or gather more measurements.
  - Keep the other variables are below the **5%** threshold.
- For categorical variables, replacing categorical variables is usually **not advisable**.
- Some common practice include replacing missing categorical variables with the **mode** of the observed ones (questionable).

# Missing Values in R

- **NA**: a missing value ("not available"), **"NA"**: a string.
- **x[1]== NA** is not a valid logical expression and will not return **FALSE** as one would expect but will return **NA**.

```
> myvector <- c(10, 20, NA, 30, 40)
> myvector
[1] 10 20 NA 30 40
> mycountry <- c("Austria", "Australia", NA, NA, "Germany", "NA")
> mycountry
[1] "Austria"    "Australia" NA           NA           "Germany"    "NA"
> is.na(myvector)
[1] FALSE FALSE TRUE FALSE FALSE
> which(is.na(myvector))
[1] 3
> x <- c(1, 4, 7, 10)
> x[4] <- NA # sets the 4th element to NA
> x
[1] 1 4 7 NA
> is.na(x) <- 1 # sets the first element to NA
> x
[1] NA 4 7 NA
```

```
> set.seed(12345)
> mydata <- matrix(round(rnorm(20), 2), ncol=5)
> mydata[sample(1:20, 3)] <- NA
> mydata
      [,1] [,2] [,3] [,4] [,5]
[1,] 0.59 0.61 NA 0.37 NA
[2,] 0.71 -1.82 -0.92 0.52 -0.33
[3,] -0.11 0.63 -0.12 -0.75 1.12
[4,] -0.45 -0.28 1.82 NA 0.30
> which(colSums(is.na(mydata)) > 0)
[1] 3 4 5
```

NOTE: **NULL** denotes something which never existed and cannot exist at all.

# NA in Summary Functions

11/70

- Most of the statistical summary functions (**mean**, **var**, **sum**, **min**, **max**, etc.) accept an argument called **na.rm**, which can be set to **TRUE** if you want missing values to be removed before the summary is calculated. (default : **FALSE**)

```
> x <- c(1, 4, NA, 10)
> summary(x)
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.    NA's
   1.0    2.5     4.0     5.0    7.0    10.0         1
> mean(x)
[1] NA
> sd(x)
[1] NA
> mean(x, na.rm=TRUE)
[1] 5
> sd(x, na.rm=TRUE)
[1] 4.582576
> x[!is.na(x)]
[1] 1 4 10
```

# NA in Modeling Functions

12/70

```
> mydata <- as.data.frame(matrix(sample(1:20, 8), ncol = 2))
> mydata[4, 2] <- NA
> names(mydata) <- c("y", "x")
> mydata
  y x
1 1 19
2 6 12
3 10 2
4 4 NA
> lm(y~x, data = mydata)

Call:
lm(formula = y ~ x, data = mydata)

Coefficients:
(Intercept)          x
    11.3927     -0.5205

> lm(y~x, data = mydata, na.action = na.omit)

Call:
lm(formula = y ~ x, data = mydata, na.action = na.omit)

Coefficients:
(Intercept)          x
    11.3927     -0.5205

> lm(y~x, data = mydata, na.action = na.fail)
Error in na.fail.default(list(y = c(1L, 6L, 10L, 4L), x = c(19L, 12L, 2L, NA)),
  missing values in object
```

# Other Special Values in R

13/70

- **NaN**: "not a number" which can arise for example when we try to compute the undeterminate  $0/0$ .

```
> x <- c(1, 0, 10)
> x/x
[1] 1 NaN 1
> is.nan(x/x)
[1] FALSE TRUE FALSE
```

- **Inf** which results from computations like  $1/0$ .
- Using the functions `is.finite()` and `is.infinite()` we can determine whether a number is finite or not.

```
> 1/x
[1] 1.0 Inf 0.1
> is.finite(1/x)
[1] TRUE FALSE TRUE
>
> -10/x
[1] -10 -Inf -1
> is.infinite(-10/x)
[1] FALSE TRUE FALSE
```

```
> exp(-Inf)
[1] 0
> 0/Inf
[1] 0
> Inf - Inf
[1] NaN
> Inf/Inf
[1] NaN
```

# R Packages for Dealing With Missing Values

14/70

- **Amelia (Amelia II)**: A Program for Missing Data
- **hot.deck**: Multiple Hot-Deck Imputation
- **HotDeckImputation**: Hot Deck Imputation Methods for Missing Data
- **impute**: (Bioconductor) Imputation for Microarray Data
- **mi**: Missing Data Imputation and Model Checking
- **mice**: **Multivariate Imputation by Chained Equations**
- **missForest**: Nonparametric Missing Value Imputation using Random Forest
- **missMDA**: Handling Missing Values with Multivariate Data Analysis (e.g., `imputePCA`, `imputeMCA`),
- **mitools**: Tools for Multiple Imputation of Missing Data
- **norm**: Analysis of Multivariate Normal Datasets with Missing Values
- **VIM**: **Visualization and Imputation of Missing Values**
- R packages support for missing values imputation.
  - **Hmisc**: Harrell Miscellaneous
  - **survey**: analysis of complex survey samples
  - **Zelig**: Everyone's Statistical Software
  - **rfImpute{randomForest}**: Imputations by randomForest
  - **imputation{rminer}**: Data Mining Classification and Regression Methods, Missing data imputation (e.g. substitution by value or hotdeck method).
  - **impute.svd{bcv}**: Cross-Validation for the SVD (Bi-Cross-Validation), Missing value imputation via a low-rank SVD approximation estimated by the EM algorithm.
  - **mlr**: Machine Learning in R provides several imputation methods.  
<https://mlr-org.github.io/mlr-tutorial/release/html/index.html>

Package "**imputation**" was removed from the CRAN. (Archived on 2014-01-14)

- **mice**: Multivariate Imputation by **Chained Equations** in R by Stef van Buuren.
- Imputing missing values on:
  - **Continuous data**: Predictive mean matching, Bayesian linear regression, Linear regression ignoring model error, Unconditional mean imputation etc.
  - **Binary data**: Logistic Regression, Logistic regression with bootstrap
  - **Categorical data** (More than 2 categories) - Polytomous logistic regression, Proportional odds model etc.
  - **Mixed data** (Can work for both Continuous and Categorical) - CART, Random Forest, Sample (Random sample from the observed values).

Source: <http://www.listendata.com/2015/08/missing-imputation-with-mice-package-in.html>



# Generates Multivariate Imputations by Chained Equations (MICE) 16/70

```
mice(data, m = 5, method = vector("character", length = ncol(data)),
      predictorMatrix = (1 - diag(1, ncol(data))),
      visitSequence = (1:ncol(data))[apply(is.na(data), 2, any)],
      form = vector("character", length = ncol(data)),
      post = vector("character", length = ncol(data)), defaultMethod = c("pmm",
      "logreg", "polyreg", "polr"), maxit = 5, diagnostics = TRUE,
      printFlag = TRUE, seed = NA, imputationMethod = NULL,
      defaultImputationMethod = NULL, data.init = NULL, ...)
```

```
> methods(mice)
```

```
[1] mice.impute.2l.norm      mice.impute.2l.pan      mice.impute.2lonly.mean
[4] mice.impute.2lonly.norm  mice.impute.2lonly.pmm  mice.impute.cart
[7] mice.impute.fastpmm      mice.impute.lda         mice.impute.logreg
```

	Method	Description	Scale type	Default
[10] mice.impute	pmm	Predictive mean matching	numeric	Y
[13] mice.impute	norm	Bayesian linear regression	numeric	
[16] mice.impute	norm.nob	Linear regression, non-Bayesian	numeric	
[19] mice.impute	mean	Unconditional mean imputation	numeric	
[22] mice.impute	2L.norm	Two-level linear model	numeric	
[25] mice.theme	logreg	Logistic regression	factor, 2 levels	Y
see '?methods' f	polyreg	Multinomial logit model	factor, >2 levels	Y
> ? mice	polr	Ordered logit model	ordered, >2 levels	Y
	lda	Linear discriminant analysis	factor	
	sample	Random sample from the observed data	any	

# Exploring Missing Data

17/70

```
> head(airquality)
  Ozone Solar.R Wind Temp Month Day
1    41     190  7.4   67     5   1
2    36     118  8.0   72     5   2
3    12     149 12.6   74     5   3
4    18     313 11.5   62     5   4
5     NA      NA 14.3   56     5   5
6    28      NA 14.9   66     5   6

> dim(airquality)
[1] 153  6

> mydata <- airquality
> mydata[4:10, 3] <- rep(NA, 7)
> mydata[1:5, 4] <- NA
>
> # Use numerical variables as examples here.
> # Ozone is the variable with the most missing datapoints.
> summary(mydata)
```

Ozone		Solar.R		Wind		Temp		Month		Day	
Min.	: 1.00	Min.	: 7.0	Min.	: 1.700	Min.	:57.00	Min.	:5.000	Min.	: 1.0
1st Qu.:	18.00	1st Qu.:	115.8	1st Qu.:	7.400	1st Qu.:	73.00	1st Qu.:	6.000	1st Qu.:	8.0
Median :	31.50	Median :	205.0	Median :	9.700	Median :	79.00	Median :	7.000	Median :	16.0
Mean :	42.13	Mean :	185.9	Mean :	9.806	Mean :	78.28	Mean :	6.993	Mean :	15.8
3rd Qu.:	63.25	3rd Qu.:	258.8	3rd Qu.:	11.500	3rd Qu.:	85.00	3rd Qu.:	8.000	3rd Qu.:	23.0
Max.	:168.00	Max.	:334.0	Max.	:20.700	Max.	:97.00	Max.	:9.000	Max.	:31.0
NA's	:37	NA's	:7	NA's	:7	NA's	:5				

Source: <http://www.r-bloggers.com/imputing-missing-data-with-r-mice-package/>

# Visualizing the Pattern of Missing Data

18/70


```
> library(mice)
> md.pattern(mydata)
      Month Day Temp Solar.R Wind Ozone
104      1   1   1      1     1     1  0
 34      1   1   1      1     1     0  1
  4      1   1   1      0     1     1  1
  3      1   1   1      1     0     1  1
  3      1   1   0      1     1     1  1
  1      1   1   1      0     1     0  2
  1      1   1   1      1     0     0  2
  1      1   1   1      0     0     1  2
  1      1   1   0      1     0     1  2
  1      1   1   0      0     0     0  4
      0   0   5      7     7    37 56
```

```
> library(VIM)
> mydata.aggrplot <- aggr(mydata,
  col=c('lightblue','red'), numbers=TRUE,
  prop = TRUE, sortVars=TRUE,
  labels=names(mydata), cex.axis=.7, gap=3)
```

Variables sorted by number of missings:

Variable	Count
Ozone	0.24183007
Solar.R	0.04575163
Wind	0.04575163
Temp	0.03267974
Month	0.00000000
Day	0.00000000

Aggregation Plot



# Matrix Plot



19/70

```
> matrixplot(mydata)
```

Click in a column to sort by the corresponding variable.

To regain use of the VIM GUI and the R console, click outside the plot region.

Matrix plot sorted by variable 'Ozone'.



# Number of Observations Per Patterns for All Pairs of Variables

20/70

V2	V	partial	complete
	x	all missing	partial
		x	v
		V1	

- **rr**: response-response, both variables are observed
- **rm**: response-missing, row observed, column missing
- **mr**: missing-response, row missing, column observed
- **mm**: missing-missing, both variables are missing

```
> md.pairs(mydata)
```

**\$rr**

	Ozone	Solar.R	Wind	Temp	Month	Day
Ozone	116	111	111	112	116	116
Solar.R	111	146	141	142	146	146
Wind	111	141	146	143	146	146
Temp	112	142	143	148	148	148
Month	116	146	146	148	153	153
Day	116	146	146	148	153	153

**\$rm**

	Ozone	Solar.R	Wind	Temp	Month	Day
Ozone	0	5	5	4	0	0
Solar.R	35	0	5	4	0	0
Wind	35	5	0	3	0	0
Temp	36	6	5	0	0	0
Month	37	7	7	5	0	0
Day	37	7	7	5	0	0

**\$mr**


	Ozone	Solar.R	Wind	Temp	Month	Day
Ozone	0	35	35	36	37	37
Solar.R	5	0	5	6	7	7
Wind	5	5	0	5	7	7
Temp	4	4	3	0	5	5
Month	0	0	0	0	0	0
Day	0	0	0	0	0	0

**\$mm**

	Ozone	Solar.R	Wind	Temp	Month	Day
Ozone	37	2	2	1	0	0
Solar.R	2	7	2	1	0	0
Wind	2	2	7	2	0	0
Temp	1	1	2	5	0	0
Month	0	0	0	0	0	0
Day	0	0	0	0	0	0

# Marginplot

```
> marginplot(mydata[,c("Ozone", "Solar.R")], col = c("blue", "red"))
```



- The **blue box** plot located on the left and bottom margins shows the distribution of the **non-missing** datapoints.
- The **red box** plot on the left shows the distribution of `Solar.R` with `Ozone` **missing**.
- If our assumption of **MCAR** data is correct, then we expect the **red** and **blue** box plots to be very similar.

# List-wise Deletion

- Also called the **complete case analysis**.
- The use of this method is only justified if the missing data generation mechanism is **MCAR**.

```
> mdata <- matrix(rnorm(15), nrow=5)
> mdata[sample(1:15, 4)] <- NA
> mdata <- as.data.frame(mdata)
> mdata
```

	V1	V2	V3
1	-0.62222501	1.0807983	NA
2	0.07124865	0.5216675	-0.08334454
3	1.70707399	0.1004917	0.88197789
4	NA	-0.6595201	-0.08387860
5	NA	1.6138847	NA

```
> (x1 <- na.omit(mdata))
```

	V1	V2	V3
2	0.07124865	0.5216675	-0.08334454
3	1.70707399	0.1004917	0.88197789

```
> (x2 <- mdata[complete.cases(mdata),])
```

	V1	V2	V3
2	0.07124865	0.5216675	-0.08334454
3	1.70707399	0.1004917	0.88197789

```
> mdata[!complete.cases(mdata),]
```

	V1	V2	V3
1	-0.622225	1.0807983	NA
4	NA	-0.6595201	-0.0838786
5	NA	1.6138847	NA

快速分析一下，得知資料大概狀況



# Pairwise Deletion

- To compute a **covariance matrix**, each **two cases** will be used for which the values of both corresponding variables **are available**.
- This can result in covariance or correlation matrices which are not positive semi-definite, as well as NA entries **if there are no complete pairs** for the given pair of variables.

```
> mdata
      V1      V2      V3
1 -0.62222501  1.0807983    NA
2  0.07124865  0.5216675 -0.08334454
3  1.70707399  0.1004917  0.88197789
4           NA -0.6595201 -0.08387860
5           NA  1.6138847    NA

> cov(mdata)
      V1      V2 V3
V1 NA      NA NA
V2 NA 0.7694197 NA
V3 NA      NA NA

> cov(mdata, use = "all.obs")
Error in cov(mdata, use = "all.obs") :
missing observations in cov/cor

> cov(mdata, use = "complete.obs")
      V1      V2      V3
V1  1.3379623 -0.34448500  0.7895494
V2 -0.3444850  0.08869452 -0.2032852
V3  0.7895494 -0.20328521  0.4659237
```

```
> cov(mdata, use = "na.or.complete")
      V1      V2      V3
V1  1.3379623 -0.34448500  0.7895494
V2 -0.3444850  0.08869452 -0.2032852
V3  0.7895494 -0.20328521  0.4659237

> cov(mdata, use = "pairwise")
      V1      V2      V3
V1  1.4304107 -0.56002326  0.78954945
V2 -0.5600233  0.76941970  0.05468712
V3  0.7895494  0.05468712  0.31078774
```

# Mean Substitution

- A very simple but popular approach is to substitute means for the missing values.
- This method produces **biased estimates** and can severely **distort the distribution** of the variable in which missing values are substituted.
- Due to these **distributional problems**, it is often **recommended to ignore missing values** rather than impute values by mean substitution (Little and Rubin, 1989.)

```
mean.subst <- function(x) {
  x[is.na(x)] <- mean(x, na.rm = TRUE)
  x
}
```

```
> mdata
      V1      V2      V3
1 -0.62222501  1.0807983    NA
2  0.07124865  0.5216675 -0.08334454
3  1.70707399  0.1004917  0.88197789
4      NA -0.6595201 -0.08387860
5      NA  1.6138847    NA
> mdata.mip <- apply(mdata, 2, mean.subst)
> mdata.mip
      V1      V2      V3
[1,] -0.62222501  1.0807983  0.23825158
[2,]  0.07124865  0.5216675 -0.08334454
[3,]  1.70707399  0.1004917  0.88197789
[4,]  0.38536588 -0.6595201 -0.08387860
[5,]  0.38536588  1.6138847  0.23825158
```

# K-Nearest Neighbour Imputation

- KNN imputation searches for the  $k$ -nearest observations (relative to the observation which has to be imputed) and replaces the missing value with the mean of the found  $k$  observations.
- It is recommended to use the (weighted) median instead of the arithmetic mean.
- **KNN minimize** data modeling assumptions and take advantage of the **correlation structure** of the data.

	$C_1$	$C_2$	$\dots$	$C_j$	$\dots$	$C_n$
$g_1$	*	✓	*	✓	✓	✓
$g_2$	✓	✓	✓	✓	✓	✓
$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$
$g_i$	✓	✓	*	✓	✓	✓
$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$
$g_m$	*	*	*	*	*	*

**KNNimpute**

**Model:**

$$\{g_{(k)}, k = 1, 2, \dots, K\} = \underset{k}{\operatorname{args\,max}} \operatorname{Corr}(g_1, g_i)$$

$$\{g_{(k)}, k = 1, 2, \dots, K\} = \underset{k}{\operatorname{args\,min}} \operatorname{Dist}(g_1, g_i)$$

$C$ : Observed  $C_i$ 's without missing values

**Imputation:**

$$\text{Average} \quad \widehat{C_1(g_1)} = \frac{1}{K} \sum_{k=1}^K C_1(g_k)$$

$$\text{Weighted Average} \quad \widehat{C_1(g_1)} = \frac{\sum_{k=1}^K w_k C_1(g_k)}{\sum_{k=1}^K w_k}$$

$$w_k = \frac{1}{\sum_{j \in C} [C_j(g_k) - C_1(g_1)]^2}$$

# k-Nearest Neighbour Imputation

## Description

k-Nearest Neighbour Imputation based on a variation of the Gower Distance for numerical, categorical, ordered and semi-continuous variables.

## Usage

```
kNN(data, variable = colnames(data), metric = NULL, k = 5,
     dist_var = colnames(data), weights = NULL, numFun = median,
     catFun = maxCat, makeNA = NULL, NAcond = NULL, impNA = TRUE,
     donorcond = NULL, mixed = vector(), mixed.constant = NULL,
     trace = FALSE, imp_var = TRUE, imp_suffix = "imp", addRandom = FALSE,
     useImputedDist = TRUE, weightDist = FALSE)
```


mean  
weightedMean

```
> names(airquality)
[1] "Ozone"    "Solar.R" "Wind"     "Temp"     "Month"    "Day"
> airquality.imp.median <- kNN(airquality[1:4], k=5)
> head(airquality.imp.median)
  Ozone Solar.R Wind Temp Ozone_imp Solar.R_imp Wind_imp Temp_imp
1    41    190  7.4   67    FALSE    FALSE    FALSE    FALSE
2    36    118  8.0   72    FALSE    FALSE    FALSE    FALSE
3    12    149 12.6   74    FALSE    FALSE    FALSE    FALSE
4    18    313 11.5   62    FALSE    FALSE    FALSE    FALSE
5    35     92 14.3   56     TRUE     TRUE    FALSE    FALSE
6    28    242 14.9   66    FALSE     TRUE    FALSE    FALSE
```

- Gower JC, 1971, A General Coefficient of Similarity and Some of Its Properties. Biometrics, 857–871.
- Alexander Kowarik and Matthias Templ, 2016, Imputation with the R Package VIM, Journal of Statistical Software, Volume 74, Issue 7.

# matrixplot、自定平均函數

```
> matrixplot(airquality[1:4], interactive = F, main="airquality")  
> matrixplot(airquality.imp.median[1:4], interactive = F, main="imputed by median")
```




## 自定平均函數

```
trim_mean <- function(x){  
  mean(x, trim = 0.1)  
}
```

```
> airquality.imp.tmean <- kNN(airquality[1:4], k=5, numFun=trim_mean)
```

# Which Imputation Method?

- KNN is the most widely-used.
- Characteristics of data that may affect choice of imputation method:
  - dimensionality.
  - percentage of values missing.
  - experimental design (time series, case/control, etc.)
  - patterns of correlation in data.
- Suggestion:
  - add (same percentage) artificial missing values to your (complete cases) data set.
  - impute them with various methods, see which is best (since you know the real value)





# Classical (Numerical) Data Table <sup>29/70</sup>

$i$ th subject  
( $i$ th sample)

$j$ th variable

UID	alpha0	alpha7	alpha14	alpha21	alpha28	alpha35	alpha42
YAR007C	-0.48	-0.42	0.87	0.92	0.67	-0.18	-0.35
YBL035C	-0.39	-0.58	1.08	1.21	0.52	-0.33	-0.58
YBR023C	0.87	0.25	-0.17	0.18	-0.13	-0.44	-0.13
YBR067C	1.57	1.03	1.22	0.31	0.16	-0.49	-1.02
YBR088C	-1.15	-0.86	1.21	1.62	1.12	0.16	-0.44
YBR278W	0.04	-0.12	0.31	0.16	0.17	-0.06	0.08
YCL055W	2.95	0.45	-0.4	-0.66	-0.59	-0.38	-0.76
YDL003W	-1.22	-0.74	1.34	1.5	0.63	0.29	-0.55
YDL055C	-0.73	-1.06	-0.79	-0.02	0.16	0.44	0.03
YDL102W	-0.58	-0.4	0.13	0.58	-0.09	0.02	-0.45
YDL164C	-0.5	-0.42	0.66	1.05	0.68	0.06	0.01
YDL197C	-0.86	-0.29	0.42	0.46	0.3	0.1	-0.63
YDL227C	-0.16	0.2877	0.17	-0.28	-0.02	-0.55	-0.04
YDR052C	-0.36	-0.03	-0.03	-0.08	-0.23	-0.25	-0.21
YDR097C	-0.72	-0.85	0.54	1.04	0.84	0.24	-0.64
YDR113C	-0.78	-0.52	0.26	0.2	0.48	0.48	0.27
YDR309C	0.6	-0.55	0.41	0.45	0.18	-0.66	-1.02
YDR356W	-0.2	-0.67	0.13	0.1	0.38	0.44	0.05
YER001W	-2.29	-0.635739	0.77	1.6	0.53	0.55	-0.38
YER070W	-1.46	-0.76	1.08	1.5	0.74	0.47	-0.7
YER095W	-0.57	0.42	1.03	1.35	0.64	0.42	-0.4
YGL163C	-0.11	0.13	0.41	0.6	0.23	0.31	0.19
YGL225W	-1.08	-0.99	-0.16	0.2	0.61	0.37	0.1
YGR109C	-1.79	0.9449	2.13	1.75	0.23	0.15	-0.66



transformation  
for each row



transformation  
for each column

transformation  
for both rows  
and columns



- to make it more closely **the assumptions** of a statistical inference procedure,
- to make it **easier to visualize** (appearance of graphs),
- to improve **interpretability**,
- to make descriptors that have been measured in **different units comparable**,
- to make the relationships among **variables linear**,
- to modify the **weights** of the variables or objects (e.g. give the same length (or norm) to all object vectors)
- to **code** categorical variables into dummy binary variables.

# 常見的資料轉換方式

31/70





```
par(mfrow=c(1,4))
raw.data <- 0:100
pa.data <- ifelse(raw.data >= 60, 1, 0)
id <- which(pa.data==1)
plot(raw.data[id], pa.data[id], main="present-absent",
+ type="l", lwd=2, col="blue", ylim=c(-1, 2), xlim=c(0, 100))
points(raw.data[-id], pa.data[-id], type="l", lwd=2, col="blue")
```

```
log.data <- log(raw.data)
plot(raw.data, log.data, main="log", type="l", lwd=2, col="blue")
```

```
sqrt10.data <- sqrt(raw.data)*10
plot(raw.data, sqrt10.data, main="sqrt*10", type="l", lwd=2, col="blue", asp=1)
abline(a=0, b=1)
```

```
trun.data <- ifelse(raw.data >= 80, 80, ifelse(raw.data < 20, 20, raw.data))
plot(raw.data, trun.data, main="truncation", type="l", lwd=2, col="blue")
```

NOTE: `apply(raw.data.matrix, 2, log)`  
`apply(raw.data.matrix, 2, function(x) sqrt(x)*10)`  
`apply(raw.data.matrix, 2, myfun)`



# 範例: Software Inspection Data

32/70



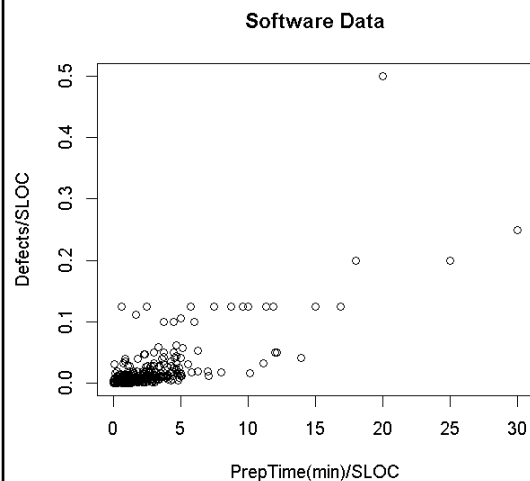
- The data were collected in response to efforts for process improvement in software testing by code inspection.
- The variables are normalized by the size of the inspection (the number of pages or **SLOC (single lines of code)**):
  - the **preparation time** in minutes (**prepage**, **prepsloc**),
  - the **total work hours** in minutes for the meeting (**mtgsloc**),
  - and the **number of defects** found (**defpage**, **defsloc**).

```
> library('R.matlab')
> data <- readMat("software.mat")
> print(data)
...
> str(data)
List of 5
 $ prepsloc: num [1:426, 1] 0.485 0.54 0.54 0.311 0.438 ...
 $ defsloc  : num [1:426, 1] 0.005 0.002 0.002 0.00328 0.00278 ...
 $ mtgsloc  : num [1:426, 1] 0.075 0.06 0.06 0.2787 0.0417 ...
 $ prepage  : num [1:491, 1] 6.15 1.47 1.47 5.06 5.06 ...
 $ defpage  : num [1:491, 1] 0.0385 0.0267 0.0133 0.0128 0.0385 ...
```

- **Interested in:** understanding the relationship between the inspection time and the number of defects found.

# 對數轉換 (Log Transformation)

33/70



```
plot(data$prepsloc, data$defsloc, xlab="PrepTime(min)/SLOC", ylab="Defects/SLOC",  
main="Software Data")
```

```
plot(log(data$prepsloc), log(data$defsloc), xlab="Log PrepTime/SLOC",  
ylab="Log Defects/SLOC", main="Software Data")
```

```
plot(log(data$prepsloc), log(data$defsloc), xlab="Log PrepTime/SLOC",  
ylab="Log Defects/SLOC", main="Software Data", asp=1)
```

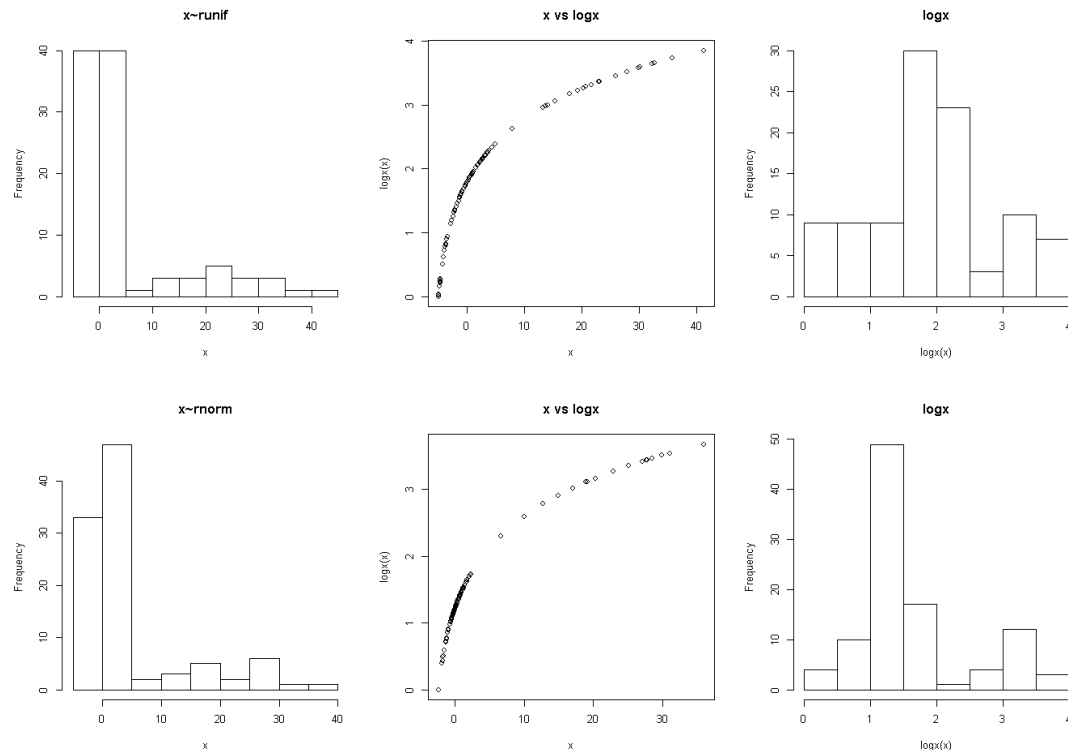
# How to Handle Negative Data Values?

## ■ Solution 1: Translate, then Transform

- $\log(x + 1 - \min(x))$

```
logx <- function(x){
  log(x + 1 - min(x))
}

x <- runif(80, min = -5, max = 5)
# x <- rnorm(80)
x <- c(x, rnorm(20, mean=20, sd=10))
par(mfrow=c(1, 3))
hist(x, main="x~runif")
plot(x, logx(x), main="x vs logx")
hist(logx(x), main="logx")
```



## ■ Solution 2: Use Missing Values

- A criticism of the previous method is that some practicing statisticians don't like to add an arbitrary constant to the data.
- They argue that a better way to handle negative values is to use missing values for the logarithm of a nonpositive number.

# Box-Cox Transformations


$$y(\lambda) = \begin{cases} \frac{y^\lambda - 1}{\lambda}, & \text{if } \lambda \neq 0; \\ \log y, & \text{if } \lambda = 0. \end{cases}$$

Box and Cox(1964)

- The aim of the Box-Cox transformations is to ensure the **usual assumptions for Linear Model hold**.

$$y \sim N(\mathbf{X}\beta, \sigma^2 \mathbf{I}_n)$$

```
x <- seq(0.5, 2, length.out=100)
bc <- function(y, lambda){
  (y^lambda - 1)/lambda
}
lambda <- seq(-2, 3, 0.5)
plot(0, 0, type="n", xlim=c(0.5, 2),
     ylim=c(-2, 2.5), main="Box-Cox transformation")
for(i in 1:length(lambda)){
  points(x, bc(x, lambda[i]), type="l", col=i)
  points(2, bc(2, lambda[i]), col=i, pch=i)
}
legend(0.7, 2.5, legend=as.character(rev(lambda)),
      lty=1, pch=length(lambda):1,
      col=length(lambda):1)
```





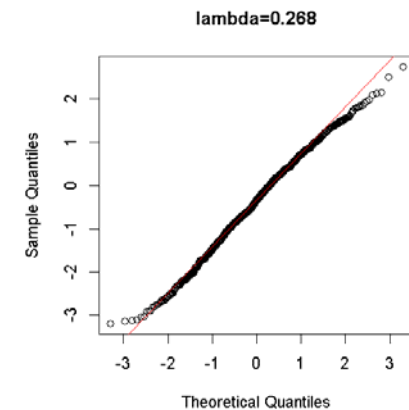



# Box-Cox Transformations

```
x <- rexp(1000)
bc <- function(y, lambda){
  (y^lambda - 1)/lambda
}
qqnorm(x); qqline(x, col="red")

bc1.x <- bc(x, 0.1)
qqnorm(bc1.x, main="lambda=0.1")
qqline(bc1.x, col="red")
bc3.x <- bc(x, 0.5)
qqnorm(bc3.x, main="lambda=0.5")
qqline(bc3.x, col="red")

bc2.x <- bc(x, 0.268)
qqnorm(bc2.x, main="lambda=0.268")
qqline(bc2.x, col="red")

hist(x, main="rexp(1000)")
hist(bc2.x, main="lamda=0.268")
```



$$\left( \Phi^{-1} \left( \frac{i - 0.5}{n} \right), x_{(i)} \right), \quad \text{for } i = 1, 2, \dots, n,$$

Source: Box-Cox Transformations: An Overview, Pengfei Li, Department of Statistics, University of Connecticut, Apr 11, 2005



# Modified Box-Cox Transformations

Manly(1971)

$$y(\lambda) = \begin{cases} \frac{e^{\lambda y} - 1}{\lambda}, & \text{if } \lambda \neq 0; \\ y, & \text{if } \lambda = 0. \end{cases}$$

Negative y's could be allowed. The transformation was reported to be successful in transform unimodal skewed distribution into normal distribution, but is not quite useful for **bimodal** or **U-shaped distribution**.

John and Draper(1980) “Modulus Transformation”

$$y(\lambda) = \begin{cases} \text{Sign}(y) \frac{(|y|+1)^\lambda - 1}{\lambda}, & \text{if } \lambda \neq 0; \\ \text{Sign}(y) \log(|y| + 1), & \text{if } \lambda = 0, \end{cases} \quad \text{Sign}(y) = \begin{cases} 1, & \text{if } y \geq 0; \\ -1, & \text{if } y < 0. \end{cases}$$

Bickel and Doksum(1981)

$$y(\lambda) = \frac{|y|^\lambda \text{Sign}(y) - 1}{\lambda}, \quad \text{for } \lambda > 0,$$

Yeo and Johnson(2000)

$$y(\lambda) = \begin{cases} \frac{(y+1)^\lambda - 1}{\lambda}, & \text{if } \lambda \neq 0, y \geq 0; \\ \log(y + 1), & \text{if } \lambda = 0, y \geq 0; \\ \frac{(1-y)^{2-\lambda} - 1}{\lambda - 2}, & \text{if } \lambda \neq 2, y < 0; \\ -\log(1 - y), & \text{if } \lambda = 2, y < 0. \end{cases}$$

Source: Box-Cox Transformations: An Overview, Pengfei Li, Department of Statistics, University of Connecticut, Apr 11, 2005

- Standardization: (called z-score): the new variate  $z$  will have a mean of zero and a variance of one. (also called centering and scaling.)

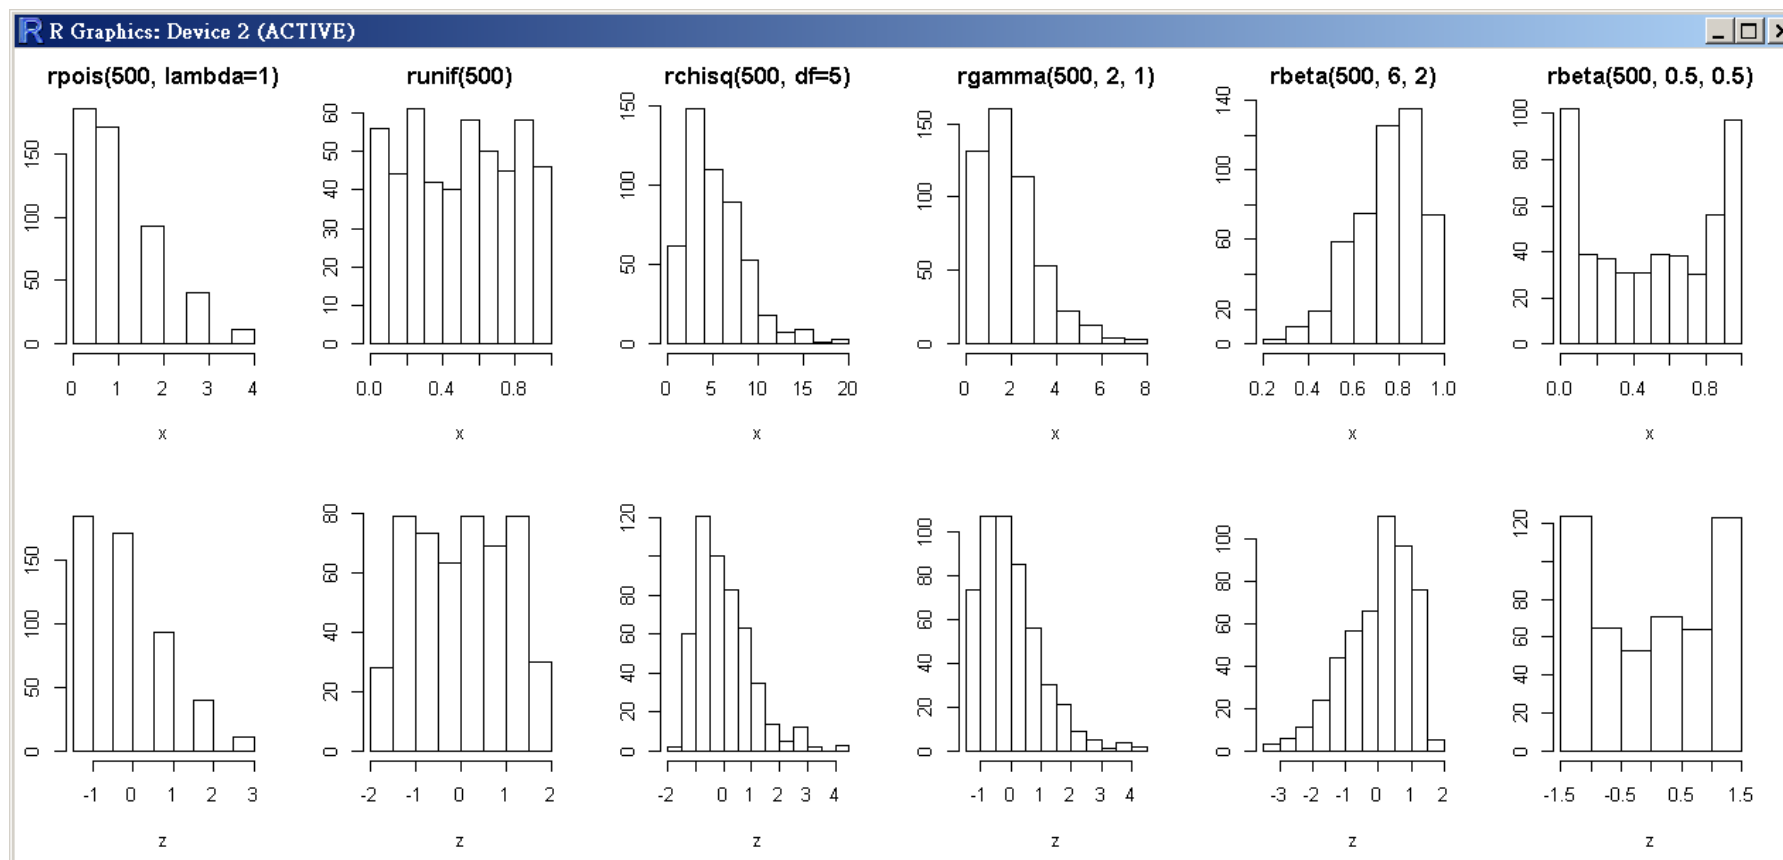
$$z_i = \frac{x_i - \bar{x}}{s}$$

- If the variables are measurements along a **different scale** or if the standard deviations for the variables are different from one another, then one variable might **dominate** the distance (or some other similar calculation) used in the analysis.
- Standardization is useful for comparing variables expressed in different units.

# 標準化 (Standardization)

39/70

Standardization makes no difference to the shape of a distribution.



```
x <- rpois(500, lambda=1)
hist(x, main="rpois(500, lambda=1)"); z <- scale(x); hist(z, main="")
```

# 範例: Standardization

40/70

`airquality` {datasets}



New York Air Quality Measurements: Daily air quality measurements in New York, May to September 1973.

A data frame with 154 observations on 6 variables.

- [1] Ozone: Ozone (**ppb**)
- [2] Solar.R: Solar R (**lang**)
- [3] Wind: Wind (**mph**)
- [4] Temp: Temperature (**degrees F**)
- [5] Month: Month (**1--12**)
- [6] Day: Day of month (**1--31**)

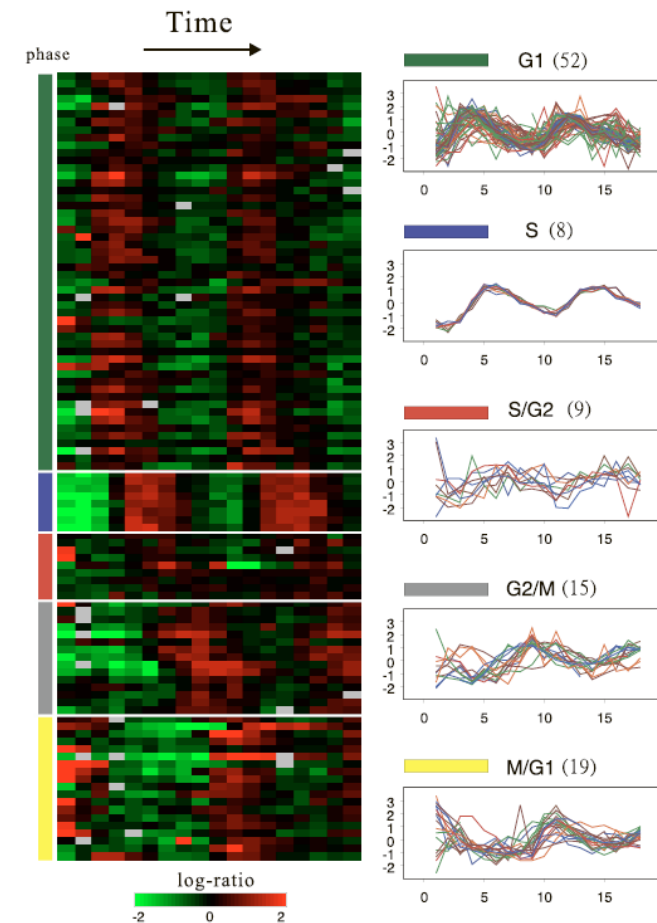

```
> head(airquality )
  Ozone Solar.R Wind Temp Month Day
1   41    190   7.4   67     5    1
2   36    118   8.0   72     5    2
3   12    149  12.6   74     5    3
4   18    313  11.5   62     5    4
5   NA     NA  14.3   56     5    5
6   28     NA  14.9   66     5    6

> r <- range(airquality[,1:4], na.rm = T)
> hist(airquality$Ozone , xlim = r)
> hist(airquality$Solar.R, xlim = r)
> hist(airquality$Wind, xlim = r)
> hist(airquality$Temp, xlim = r)
>
> airquality.std <- as.data.frame(
  apply(airquality, 2, scale))
> r.std <- c(-3, 3)
> hist(airquality.std$Ozone, xlim = r.std)
> hist(airquality.std$Solar.R, xlim = r.std)
> hist(airquality.std$Wind, xlim = r.std)
> hist(airquality.std$Temp, xlim = r.std)
```



# 範例: Microarray Data of Yeast Cell Cycle

- Lu and Wu (2010)
  - Time course data: every 7 minutes and totally 18 time points.
  - Known genes: there are 103 cell cycle-regulated genes by traditional method in G1, S, S/G2, G2/M, or M/G1. (Remove NA's: 79.)



*See also:* Using R to draw a Heatmap from Microarray Data

[http://www2.warwick.ac.uk/fac/sci/moac/people/students/peter\\_cock/r/heatmap/](http://www2.warwick.ac.uk/fac/sci/moac/people/students/peter_cock/r/heatmap/)


# Standardization in Time Series

## Microarray Gene Expression Experiments


42/70

```
cell.raw <- read.table("trad_alpha103.txt", row.names=1, header=T)
head(cell.raw)
cell.xdata <- t(scale(t(cell.raw[,2:19]), center=T, scale=T))
y.C <- as.integer(cell.raw[,1])
table(y.C)
no.cluster <- length(unique(y.C))
cellcycle.color <- c("darkgreen", "blue", "red", "gray50", "orange")
p <- ncol(cell.raw) - 1
ycolors <- cellcycle.color[y.C+1]
my.pch <- c(1:no.cluster)[y.C+1]
phase <- c("G1", "S", "S/G2", "G2/M", "M/G1")
matplot(t(cell.xdata), pch = 1:p, lty=1, type = "l", ylab="gene expression",
        col=ycolors, xlab="time", main="Time series", xaxt="n")
time.label <- parse(text=paste("t[",0:p,"]",sep=""))
axis(1, 1:(p+1), time.label)
legend("bottom", legend=phase, col=cellcycle.color, lty=1, horiz = T, lwd=2)
```

Time series



Time series



The data map for 103 cell cycle-regulated genes and the plots of time courses for each phase. Each expression profile is normalized as mean equals zero and variance 1.



# 範例: Crab Data

43/70

## crabs {MASS}

Morphological Measurements on Leptograpsus Crabs

Description: The crabs data frame has **200 rows** and **8 columns**, describing 5 morphological measurements on **50 crabs each of two colour forms and both sexes**, of the species *Leptograpsus variegatus* (紫岩蟹) collected at Fremantle, W. Australia.

**This data frame contains the following columns:**

**sp:** species - "B" or "O" for blue or orange.

**sex:** "M" or "F" for male or female

**index:** 1:50 within each of the four groups.

**FL:** carapace frontal lobe (lip) size (mm).

**RW:** carapace rear width (mm).


**CL:** carapace length (mm).

**CW:** carapace width (mm).

**BD:** body depth (mm).

```
> library(MASS)
```

```
> data(crabs)
```



Aust. J. Zool. 1974, 22, 417-25

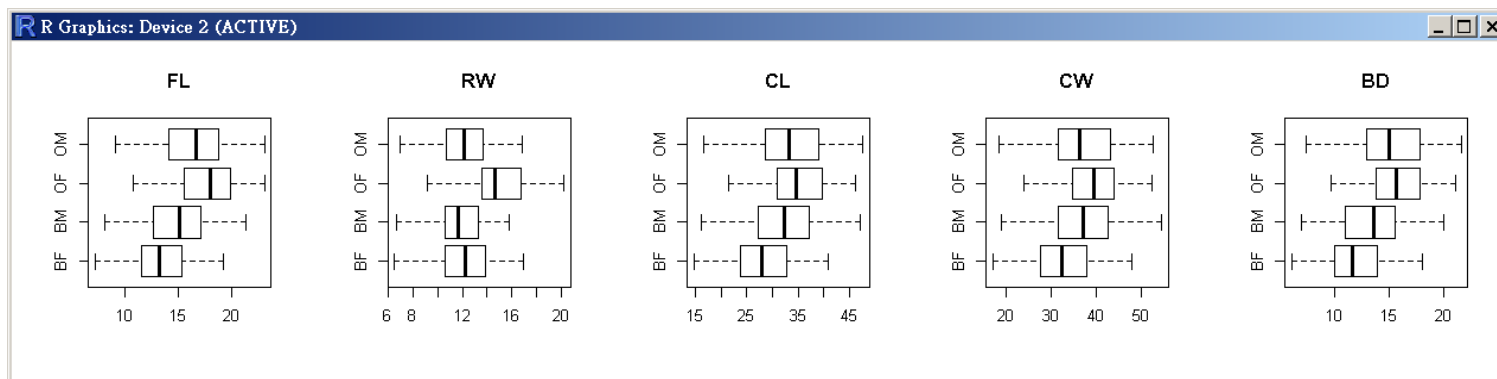




<http://www.qm.qld.gov.au/Find+out+about/Animals+of+Queensland/Crustaceans/Common+marine+crustaceans/Crabs/Purple+Swift-footed+Shore+Crab#.VhPWYiurFhs>

# 範例: Crab Data

44/70

```
boxplot(crabs$FL~crabs$sp, main="FL", horizontal=T)
```





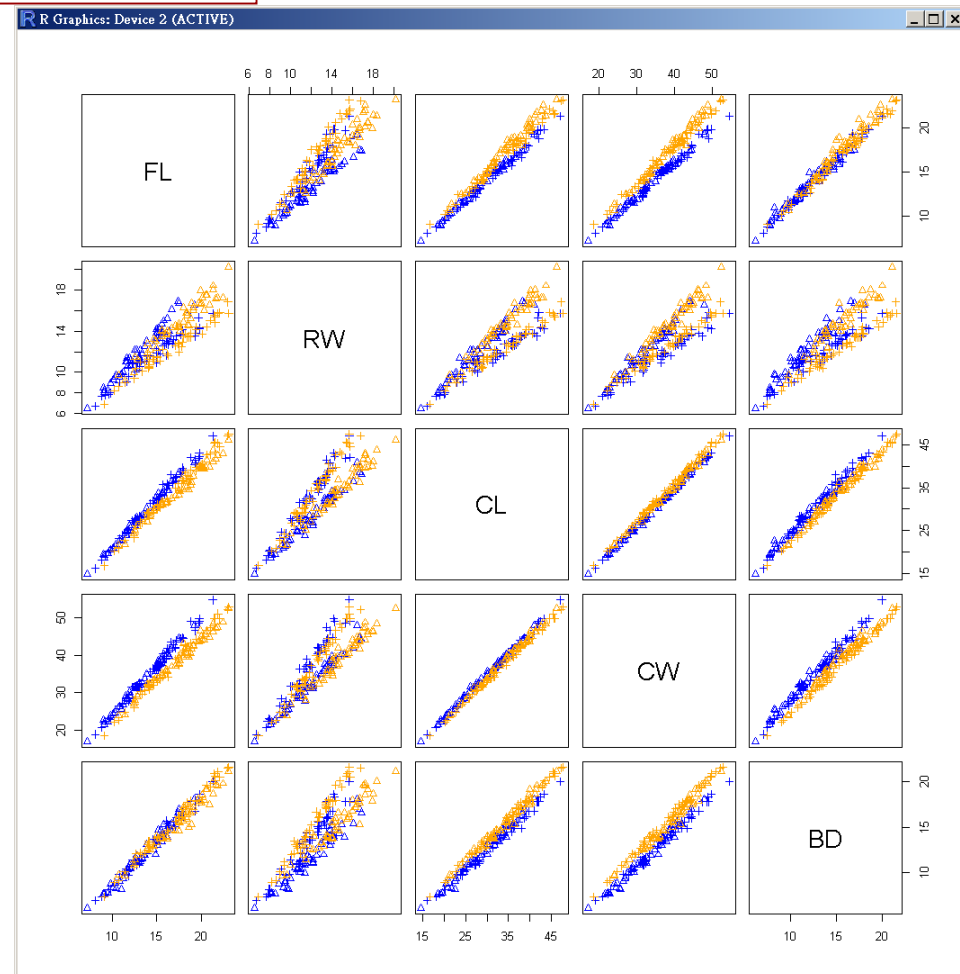
# 範例: Crab Data

45/70

```
# tri: F, cross: M  
pairs(crabs[,4:8],  
pch=as.integer(crabs$sex)+1,  
col=c("blue","orange")[as.integer(crabs$sp)])
```

- The analysis of ratios of body measurements is deeply ingrained in the taxonomic literature.
- Whether for plants or animals, certain ratios are commonly indicated in identification keys, diagnoses, and descriptions.



(Hannes Baur and Christoph Leuenberger, Analysis of Ratios in Multivariate Morphometry, Systematic Biology 60(6), 813-825.)



- The use of ratios of measurements (i.e., of body proportions), has a long tradition and is deeply ingrained in morphometric taxonomy.

Three size vectors have been commonly proposed in the literature:

- (a) isometric size  
(the arithmetic mean of  $x$ ),
- (b) allometric size,
- (c) shape-uncorrelated size.




```
par(mfrow=c(1,2))
mp <- as.integer(crabs$sex)+1
mc <- c("blue","orange")[as.integer(crabs$sp)]
isometric.size <- apply(crabs[,4:8], 1, mean)
plot(isometric.size, log(crabs$BD/crabs$RW), pch=mp, col=mc)
plot(isometric.size, log(crabs$CL/crabs$CW), pch=mp, col=mc)
```

# 範例: cDNA Microarray Gene Expression Data

47/70

微陣列資料統計分析 Statistical Microarray Data Analysis


<http://www.hmwu.idv.tw/index.php/mada>



## Why Normalization?

**Non-biological factor** can contribute to the variability of data, in order to reliably compare data from **multiple probe arrays**, differences of non-biological origin must be minimized.  
(Remove the systematic bias in the data).

- Within-Array Normalization
- Between-Array Normalization
- Paired-slides Normalization
- ...



# 要使用哪一種資料轉換方式？

48/70

- Use a transformation that other researchers **commonly use in your field**.
- Guidance for how data should be transformed, or whether a transformation should be applied at all, should come from the particular statistical analysis to be performed.
- The main criteria in choosing a transformation:
  - what works with the data?
  - what makes physical (biological, economic, whatever) sense.
- If you have a **large** number of observations, compare the effects of different transformations on the **normality** and **the homoscedasticity** of the variable.


<http://www.biostathandbook.com/transformation.html>

# Classification: k-fold Cross-Validation Error Rates


49/70

Microarray Data of  
Yeast Cell Cycle


Y X



Training Set



Test Set




2/3  
3-fold Cross-Validation  
1/3

*Possible to*

1. classification for genes
2. classification for samples (arrays)

Construct

Classification rule



Assign  
class  
labels



Assign  
class  
labels



prediction error

classification error



# Split Data into Test and Train Set According to Group Labels

50/70

```
library(caTools) # Tools: moving window statistics, GIF, Base64, ROC AUC, etc
set.seed(12345)
id <- sample.split(1:nrow(iris), SplitRatio = 0.90)
iris.train <- subset(iris, id == TRUE)
iris.test <- subset(iris, id == FALSE)
```

```
> require(caTools)
> Y <- iris[,5] # extract labels from the data
> msk <- sample.split(Y, SplitRatio=4/5)
> msk
[1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE FALSE FALSE TRUE
...
[144] TRUE TRUE TRUE FALSE TRUE TRUE FALSE
> table(Y, msk)
      msk
Y      FALSE TRUE
setosa      10   40
versicolor  10   40
virginica   10   40
> iris.train <- iris[msk, ]
> iris.test <- iris[!msk, ]
> dim(iris.train)
[1] 120  5
> dim(iris.test)
[1] 30  5
```

```
> library(caret)
> createFolds(iris$Species, k=3)
$Fold1
[1]  2  8 15 22 25 27 30 ...

$Fold2
[1]  5  6  9 10 11 12 17 ...

$Fold3
[1]  1  3  4  7 13 14 16 20...
```

```
library(caret)
id <- createDataPartition(y=iris$Species, p=0.9, list=FALSE)
iris.train <- iris[id, ]
iris.test <- iris[-id, ]
```

# Jackknife Resampling: Leave-one-out

51/70

$\hat{\theta}$  the calculated estimator of the parameter based on all  $n$  observations


$\hat{\theta}_{(.)} = \frac{1}{n} \sum_{i=1}^n \hat{\theta}_{(i)}$  the average of these "leave-one-out" estimates

$\hat{\theta}_{\text{Jack}} = n\hat{\theta} - (n-1)\hat{\theta}_{(.)}$  the resulting bias-corrected jackknife estimate


```
> # install.packages("bootstrap")
> library(bootstrap)
> x <- rnorm(20)
> theta <- function(x){mean(x)}
> (theta.hat <- theta(x))
[1] -0.1135763
> results <- jackknife(x,theta)
> results
$jack.se
[1] 0.264117

$jack.bias
[1] 2.63678e-16

$jack.values
[1] -0.091950484 -0.193139320 -0.153668397 ...
...
$call
jackknife(x = x, theta = theta)
```



```
> theta.hat.loo <- mean(results$jack.values)
> (theta.hat.jack <- n * theta.hat - (n-1) * theta.hat.loo)
[1] -0.1135763
> plot(results$jack.values, main="jackknife")
```







# 自助法、拔靴法 Bootstrap Methods

52/70


- Bootstrapping is a statistical method for estimating the **sampling distribution of an estimator** by **sampling with replacement** from the original sample, of the same size as the original sample.



Bradley Efron 1938~  
Department of Statistics, Stanford University



1993



2016

Efron : Bootstrap Methods: Another Look at the Jackknife - Project Euclid

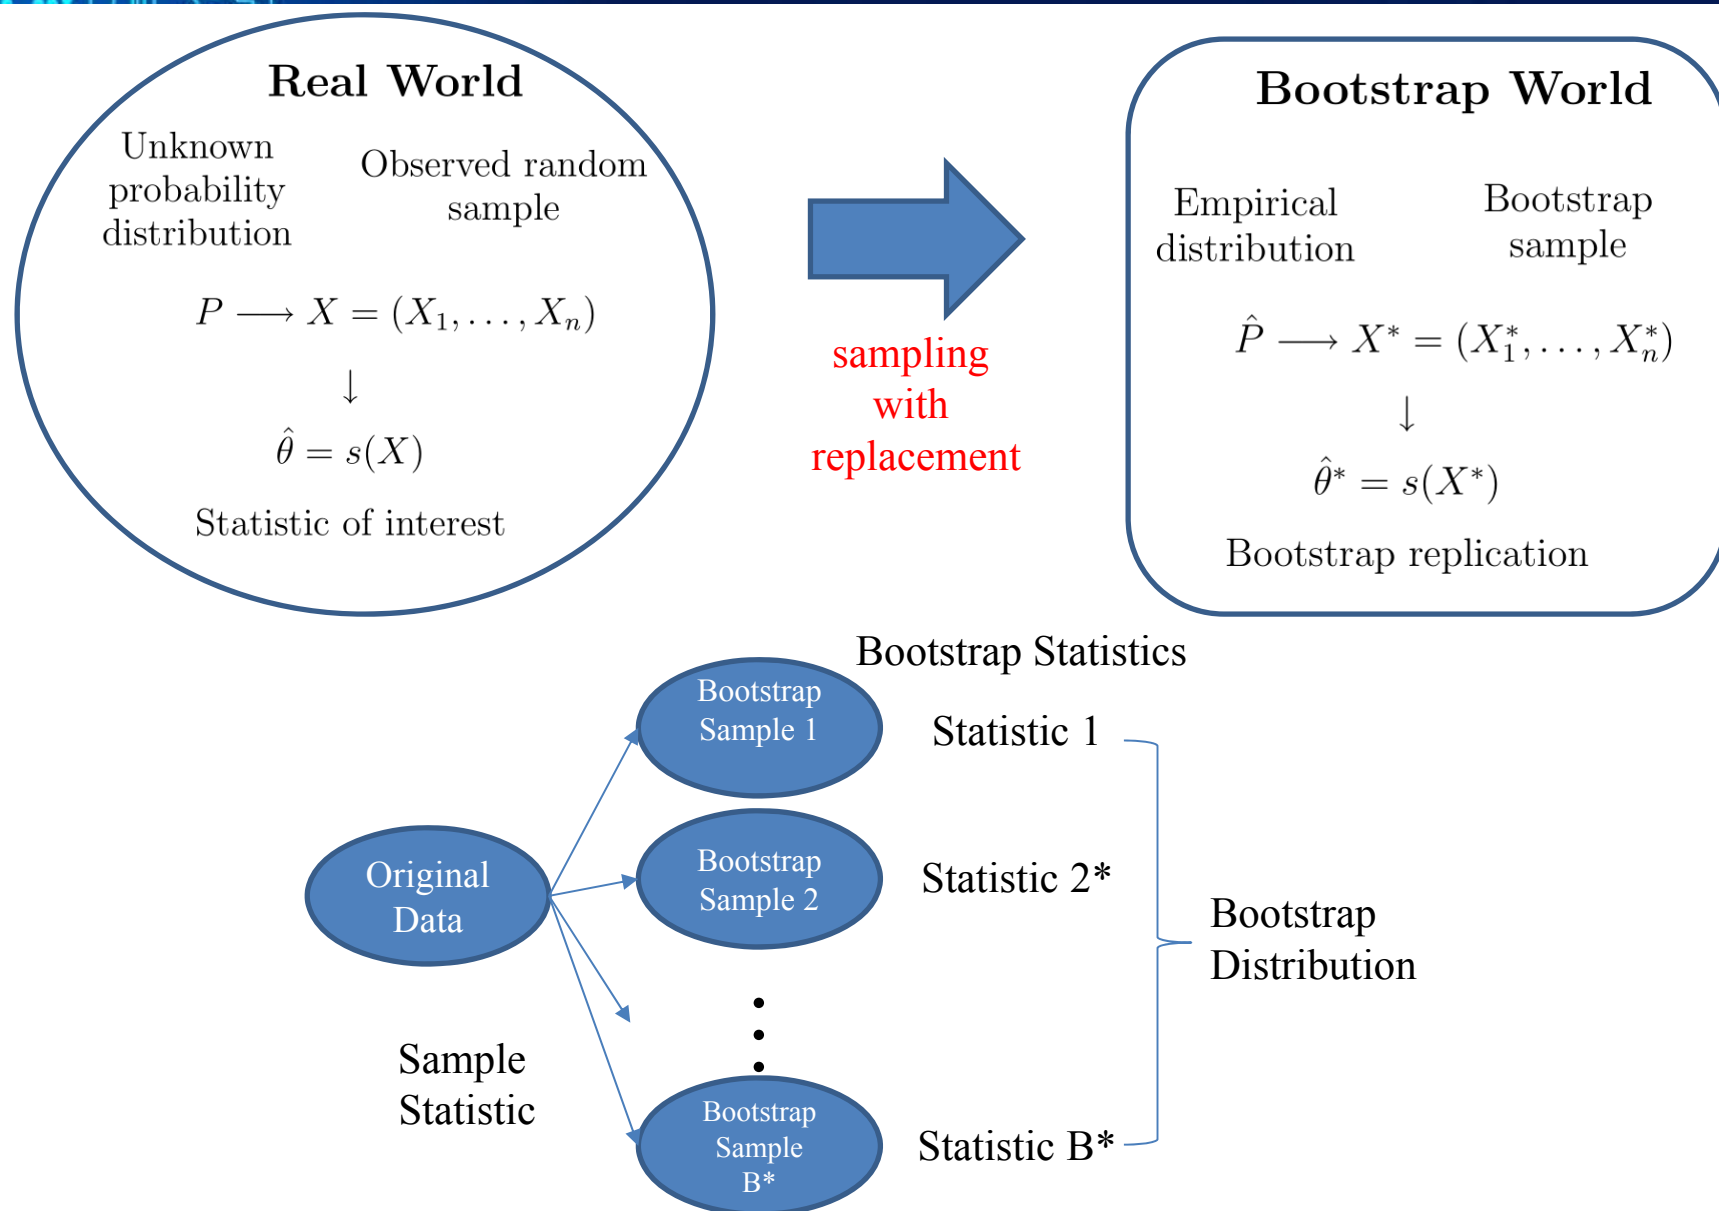
<https://projecteuclid.org/euclid.aos/1176344552> ▼ 翻譯這個網頁

由 B Efron 著作 - 1979 - 被引用 16424 次 - 相關文章

The Annals of Statistics ... Bootstrap Methods: Another Look at the Jackknife ... The jackknife is shown to be a linear approximation method for the bootstrap.



# Bootstrapping



# bootstrap Package

## Bootstrap Estimation of the Sample Mean

語法:

`bootstrap(x, nboot, theta, ..., func=NULL)`

**x:** a vector containing the data.

**nboot:** the number of bootstrap samples.

**theta:** function to be bootstrapped.

```
> # install.packages("bootstrap")
> library(bootstrap)
> set.seed(12345)
> x <- rnorm(20)
> mean(x)
[1] 0.07651681
> (x.bootstrap.mean <- bootstrap(x, 50, theta=mean))
$thetastar
 [1]  0.486197466 -0.160488357  0.274920990  0.398499864 -0.399967845  0.116086370
...
[43] -0.348643786  0.185330636 -0.070823890  0.057609481  0.062067504  0.043716794
[49] -0.279597885  0.243843620

$func.thetastar
NULL

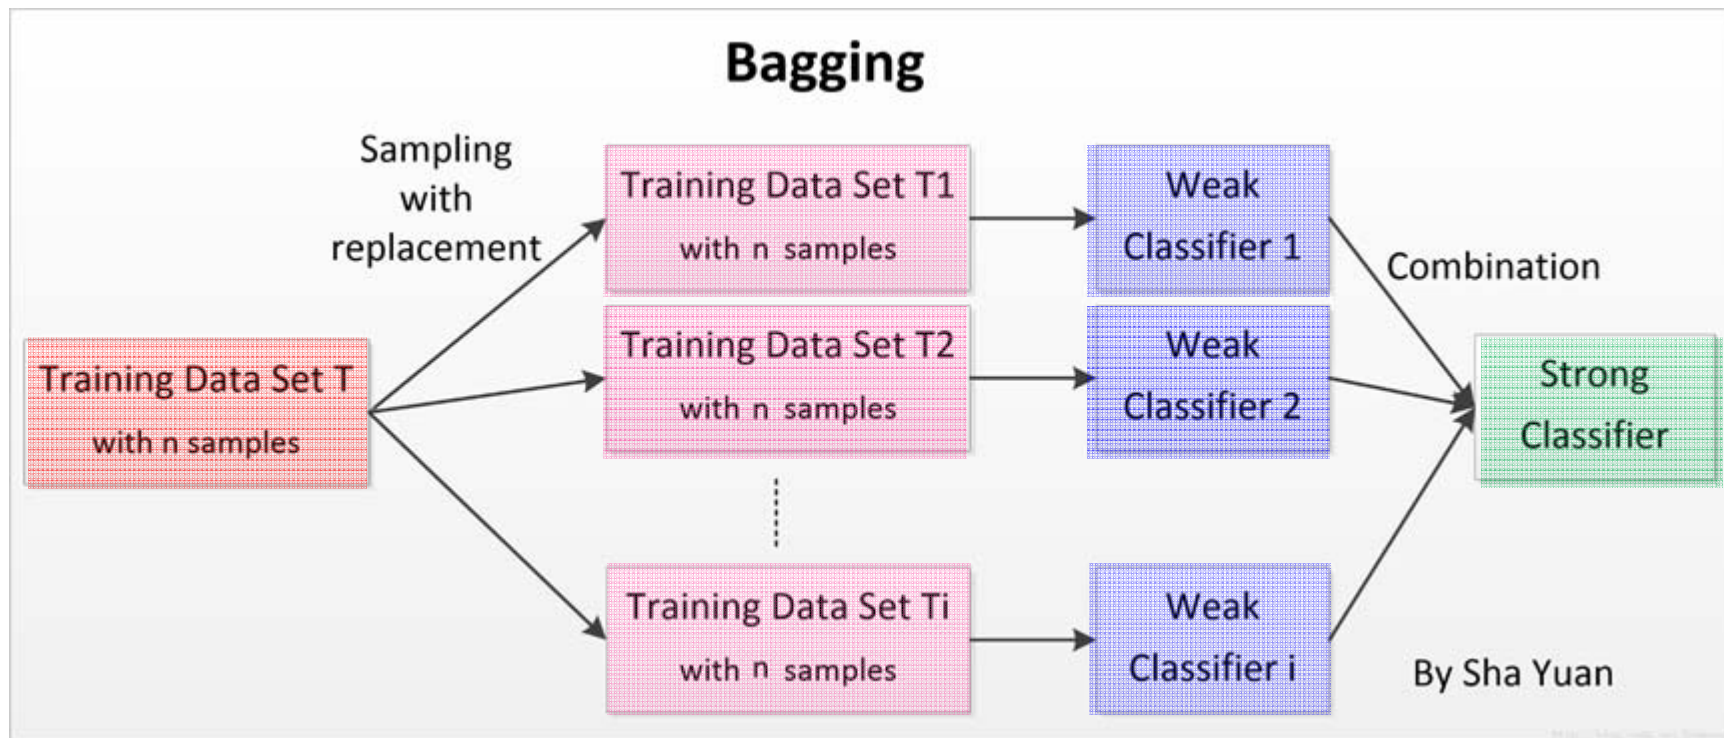
$jack.boot.val
NULL

$jack.boot.se
NULL

$call
bootstrap(x = x, nboot = 50, theta = mean)
> mean(x.bootstrap.mean$thetastar)
[1] 0.08647268
```

# Bagging: Bootstrap Aggregating

55/70



<http://blog.csdn.net/bymaymay/article/details/77824574>

- Breiman, L. (1996). Bagging predictors, Machine Learning, Vol. 26, pp. 123-140.
- Freund, Y. and Schapire, R. E. (1996). Experiments with a new boosting algorithm, Proceedings of the Thirteenth International Conference, Machine Learning.

# Boosting

$$W_{x_i}^{(1)} = N^{-1} \text{ for all } x_i.$$

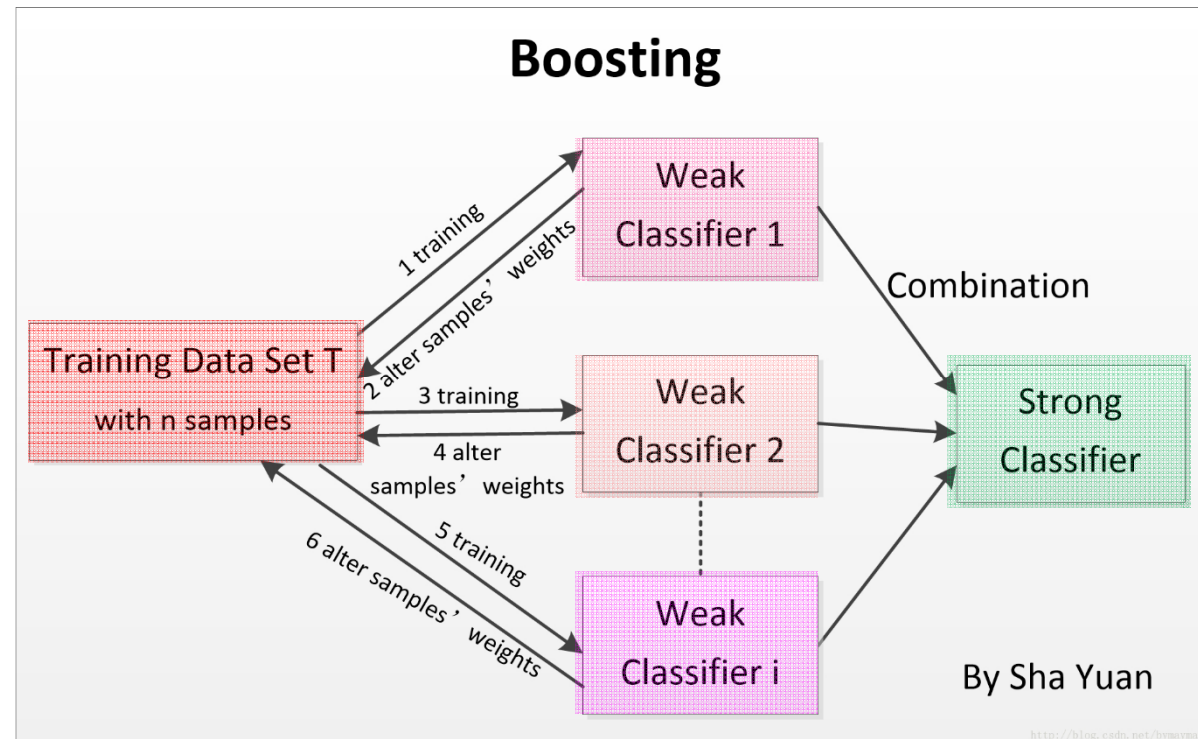
a bootstrap sample  $\mathcal{L}_t^{(B)}$   
error  $\epsilon_t$  of classifier  $\varphi_t(\mathbf{x})$

$$\epsilon_t = \sum_{\{i: \varphi_t(x_i) \neq y_i\}} W_{x_i}^{(t)}.$$

$$\beta_t = (1 - \epsilon_t) \epsilon_t^{-1}$$

$$W_{x_i}^{(t+1)} = \frac{W_{x_i}^{(t)} \beta_t^{d(i)}}{\sum_i W_{x_i}^{(t)} \beta_t^{d(i)}},$$

boosted classifier



<http://blog.csdn.net/bymaymay/article/details/77824574>

$d(i) = 1$  if  $i$ th case is classified incorrectly,  
 $d(i) = 0$ , otherwise

$$\varphi_B(x_i) = \arg \max_j \sum_{t=1}^T \log \beta_t I[\varphi_t(x_i) = j]$$

Ad-Boost.M1 (Freund and Schapire, 1996)

# Example: Apply **rpart** to Vehicle Data

```
> library(rpart); library(mlbench); library(adabag)
> data(Vehicle)
> dim(Vehicle)
[1] 846 19
> head(Vehicle)
  Comp Circ D.Circ Rad.Ra Pr.Axis.Ra Max.L.Ra Scat.Ra Elong Pr.Axis.Rect Max.L.Rect Sc.Var.Maxis
1   95   48    83   178         72    10   162   42         20        159        176
...
  Sc.Var.maxis Ra.Gyr Skew.Maxis Skew.maxis Kurt.maxis Kurt.Maxis Holl.Ra Class
1           379   184         70         6       16        187   197   van
...
6           957   264         85         5         9        181   183   bus
> table(Vehicle$Class)
bus opel saab van
218 212 217 199
```

```
> n <- nrow(Vehicle)
> sub <- sample(1:n, 2*n/3)
> Vehicle.train <- Vehicle[sub, ]
> Vehicle.test <- Vehicle[-sub, ]
```

```
> mfinal <- 10 # Defaults to mfinal=100 iterations
> maxdepth <- 5
> Vehicle.rpart <- rpart(Class ~ ., data = Vehicle.train, maxdepth = maxdepth)
> Vehicle.rpart.pred <- predict(Vehicle.rpart, newdata = Vehicle.test, type = "class")
> (tb <- table(Vehicle.rpart.pred, Observed.Class=Vehicle.test$Class))
      Observed.Class
Vehicle.rpart.pred bus opel saab van
      bus      69    10     6     2
      opel      1    25    13     3
      saab      1    34    37     8
      van       2     7     5    59
> (error.rpart <- 1 - (sum(diag(tb)) / sum(tb)))
[1] 0.3262411
```



# adabag: An R Package for Classification with 58/70 Boosting and Bagging

```
> library(adabag)
> Vehicle.adaboost <- boosting(Class ~., data = Vehicle.train, mfinal = mfinal,
+                             control = rpart.control(maxdepth=maxdepth))
> Vehicle.adaboost.pred <- predict.boosting(Vehicle.adaboost, newdata = Vehicle.test)
> Vehicle.adaboost.pred$confusion
```

	Observed Class			
Predicted Class	bus	opel	saab	van
bus	69	2	3	1
opel	1	30	16	4
saab	1	38	39	1
van	2	6	3	66

```
> Vehicle.adaboost.pred$error
[1] 0.2765957
> importanceplot(Vehicle.adaboost)
>
> # comparing error evolution in training and test set
> evol.train <- erreval(Vehicle.adaboost, newdata = Vehicle.train)
> evol.test <- erreval(Vehicle.adaboost, newdata = Vehicle.test)
> plot.erreval(evol.test, evol.train)
```

```
> sort(Vehicle.adaboost$importance, dec=T)[1:5]
      Max.L.Ra Sc.Var.maxis      Max.L.Rect
      29.623783      11.473254      9.956137
      Elong          Comp
      7.570798      6.656360
```



Alfaro, E., Gamez, M. and Garcia, N. (2013): "adabag: An R Package for Classification with Boosting and Bagging". Journal of Statistical Software, 54(2), 1–35.




# Example: 10-fold CV adaboost.M1

```
> # 10-fold CV adaboost.M1
> Vehicle.boost.cv <- boosting.cv(Class ~., data = Vehicle, v = 10, mfinal = 5,
+                               control = rpart.control(maxdepth = maxdepth))
> Vehicle.boost.cv$confusion
      Observed Class
Predicted Class bus opel saab van
      bus    209     9    11     3
      opel     1    101    72     2
      saab     0     88   117     6
      van      8     14    17   188
> Vehicle.boost.cv$error
[1] 0.2730496
```

```
> Vehicle.bag.cv <- bagging.cv(Class ~., data = Vehicle, v = 10, mfinal = 5,
+                              control = rpart.control(maxdepth = maxdepth))
> Vehicle.bag.cv$confusion
      Observed Class
Predicted Class bus opel saab van
      bus    208     20    24     2
      opel     1     91    53     1
      saab     1     81   116     5
      van      8     20    24   191
> Vehicle.bag.cv$error
[1] 0.2836879
```

## The Imbalanced Data Problem

- A dataset is said to be **unbalanced** when the class of interest (**minority class**) is much rarer than normal behaviour (**majority class**).
- **Example:** 5% of the target class represents fraudulent transactions, 95% of the target class represents legitimate transactions.
- Most learning systems are not prepared to cope with unbalanced data and several techniques have been proposed.



<http://www.srutisj.in/blog/research/statisticalmodeling/balancing-techniques-for-unbalanced-datasets-in-python-r/>



## Racing for Unbalanced Methods Selection

Re-balance or remove noisy instances in unbalanced datasets.

```
ubBalance {unbalanced}
```

### Usage

```
ubBalance(X, Y, type="ubSMOTE", positive=1,  
          percOver=200, percUnder=200,  
          k=5, perc=50, method="percPos", w=NULL, verbose=FALSE)
```

### Arguments

**X**: the input variables of the unbalanced dataset.

**Y**: the response variable of the unbalanced dataset.

**type**: the balancing technique to use (**ubOver**, **ubUnder**, **ubSMOTE**, **ubOSS**, **ubCNN**, **ubENN**, **ubNCL**, **ubTomek**).

**positive**: the majority class of the response variable.

**percOver**: parameter used in **ubSMOTE**

**percUnder**: parameter used in **ubSMOTE**

**k**: parameter used in **ubOver**, **ubSMOTE**, **ubCNN**, **ubENN**, **ubNCL**

**perc**: parameter used in **ubUnder**

**method**: parameter used in **ubUnder**

**w**: parameter used in **ubUnder**

**verbose**: print extra information (TRUE/FALSE)

**ubSMOTE {unbalanced}**: synthetic minority over-sampling technique

### Usage

```
ubSMOTE(X, Y, perc.over = 200, k = 5, perc.under = 200, verbose = TRUE)
```

**Other R packages**: **imbalanced**: Preprocessing Algorithms for Imbalanced Datasets, Imbalanced Classification in R: **ROSE** (Random Over Sampling Examples) and **DMwR** (Data Mining with R).

# The Balancing Technique


- **ubOver**: replicates randomly some instances from the minority class in order to obtain a final dataset with the same number of instances from the two classes.
- **ubUnder**: removes randomly some instances from the majority (negative) class and keeps all instances in the minority (positive) class in order to obtain a more balanced dataset.
- **ubCNN**: **Condensed Nearest Neighbor** selects the subset of instances that are able to correctly classifying the original datasets using a **one-nearest neighbor** rule.
- **ubENN**: **Edited Nearest Neighbor** removes any example whose class label differs from the class of at least two of its three nearest neighbors.
- **ubNCL**: **Neighborhood Cleaning Rule** modifies the Edited Nearest Neighbor method by increasing the role of data cleaning.
  - Firstly, NCR **removes** negatives examples which are misclassified by their 3-nearest neighbors.
  - Secondly, the neighbors of each positive examples are found and the ones belonging to the majority class are **removed**.

# The Balancing Technique

63/70

- **ubTomek**: finds the points in the dataset that are **tomek link** using 1-NN and then removes only majority class instances that are tomek links.

x's nearest neighbor is y  
y's nearest neighbor is x  
x and y are different classes




[http://contrib.scikit-learn.org/imbalanced-learn/stable/auto\\_examples/under-sampling/plot\\_illustration\\_tomek\\_links.html](http://contrib.scikit-learn.org/imbalanced-learn/stable/auto_examples/under-sampling/plot_illustration_tomek_links.html)

- **ubOSS: One Side Selection** is an undersampling method resulting from the application of **Tomek links** followed by the application of **Condensed Nearest Neighbor**.

# The Balancing Technique

- **ubSMOTE**: **synthetic minority over-sampling technique** generates new examples by filling empty areas among the positive instances.



SMOTE: Synthetic Minority Over-sampling Technique | Journal of ...

<https://jair.org/papers/paper953.html> ▼ 翻譯這個網頁

由 NV Chawla 著作 - 2002 - 被引用 5757 次 - 相關文章

An approach to the construction of classifiers from imbalanced datasets is described. A dataset is imbalanced if the classification categories are not ...

# Ionosphere (電離層) dataset

## ubIonosphere {unbalanced}

65/70

- The datasets is a modification of Ionosphere dataset contained in "mlbench" package.

```
> # install.packages("unbalanced")
> library(unbalanced)
> p <- ncol(ubIonosphere)
> y <- ubIonosphere$Class
> x <- ubIonosphere[, -p]
> data <- ubBalance(X=x, Y=y, type="ubOver", k=0)
> overData <- data.frame(data$X, Class=data$Y)
> table(overData$Class)
 0    1
225 225
```

perc: percentage of sampling

```
> data <- ubBalance(X=x, Y=y, type="ubUnder", perc=50, method="percPos")
> underData <- data.frame(data$X, Class=data$Y)
> table(underData$Class)
 0    1
126 126
```

```
> bdata <- ubBalance(X=x, Y=y, type="ubSMOTE", percOver=300, percUnder=150, verbose=TRUE)
Proportion of positives after ubSMOTE : 47.06 % of 1071 observations
```

```
> str(bdata)
```

List of 3

```
$ X      : 'data.frame': 1071 obs. of 32 variables:
 ..$ V3 : num [1:1071] -0.787 1 1 0.5 1 ...
...
..$ V34: num [1:1071] -0.576 0.714 -0.243 0.174 -0.892 ...
$ Y      : Factor w/ 2 levels "0","1": 2 1 1 1 1 2 1 2 1 2 ...
$ id.rm: logi NA
> table(bdata$Y)
```

```
 0    1
567 504
```

per.over/100 : number of new instances generated for each rare instance

perc.under/100: number of "normal" (majority class) instances that are randomly selected for each smoted observation.

```
> data(ubIonosphere)
> dim(ubIonosphere)
[1] 351 33
> head(ubIonosphere)
      V3      V4      V34 Class
1 0.99539 -0.05889 ... -0.45300      0
...
6 0.02337 -0.00592 ...  0.12011      1
> table(ubIonosphere$Class)
 0    1
225 126
```

K=0: sample with replacement from the minority class until we have the same number of instances in each class.

If K>0: sample with replacement from the minority class until we have k-times the original number of minority instances

# Compare the Performances using SVM

66/70

```
> set.seed(12345)
> n <- nrow(ubIonosphere) # 351
> no.train <- floor(0.5*n) # 175, keep half for training and half for testing
> id <- sample(1:n, no.train)
> x.train <- x[id, ] # 175 x 32
> y.train <- y[id]
> x.test <- x[-id, ] # 176 32
> y.test <- y[-id]
>
> library(e1071)
> model1 <- svm(x.train, y.train)
> y.pred1 <- predict(model1, x.test)
> table(y.pred1, y.test)
      y.test
y.pred1  0   1
      0 113  10
      1   4  49
>
> # rebalance the training set before building a model
> balancedData <- ubBalance(X=x.train, Y=y.train, type="ubSMOTE",
                           percOver=200, percUnder=150)
> table(balancedData$Y)
  0   1
201 201
```

```
> model2 <- svm(balancedData$X, balancedData$Y)
> y.pred2 <- predict(model2, x.test)
> table(y.pred2, y.test)
      y.test
y.pred2  0   1
      0 112   8
      1   5  51
```

# ubRacing {unbalanced}

## Racing for Strategy Selection

67/70

Selecting the best technique to re-balance or remove noisy instances in unbalanced datasets.

```
ubRacing(formula, data, algo, positive=1, ncore=1, nFold=10, maxFold=10, maxExp=100,  
         stat.test="friedman", metric="f1", ubConf, verbose=FALSE, ...)
```

### Arguments

**algo**: the classification algorithm to use with the mlr package.

**positive**: label of the positive (minority) class.

**ncore**: the number of core (parallel execution) to use in the Race.

```
> set.seed(1234)  
> # load(url("http://www.ulb.ac.be/di/map/adalpozz/data/creditcard.Rdata"))  
> load("creditcard.Rdata")  
> str(creditcard)  
'data.frame': 284807 obs. of 31 variables:  
 $ Time : num 0 0 1 1 2 2 4 7 7 9 ...  
 $ V1 : num -1.36 1.192 -1.358 -0.966 -1.158 ...  
 ...  
 $ V28 : num -0.0211 0.0147 -0.0598 0.0615 0.2152 ...  
 $ Amount: num 149.62 2.69 378.66 123.5 69.99 ...  
 $ Class : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...  
> table(creditcard$Class)  
  
 0      1  
284315 492  
> # configuration of the sampling method used in the race  
> ubConf <- list(percOver=200, percUnder=200, k=2, perc=50, method="percPos", w=NULL)  
> # Race with 5 trees in the Random Forest  
> results <- ubRacing(Class ~., creditcard, "randomForest",  
+                   positive=1, metric="auc", ubConf=ubConf, ntree=5)
```

The function **ubRacing** compares the 8 unbalanced methods (ubUnder, ubOver, ubSMOTE, ubOSS, ubCNN, ubENN, ubNCL, ubTomek) against the unbalanced distribution.

# Racing for Strategy Selection

Racing for unbalanced methods selection in 10 fold CV

Number of candidates.....9  
 Max number of folds in the CV.....10  
 Max number of experiments.....100  
 Statistical test.....Friedman test

Markers:

- x No test is performed.
- The test is performed and some candidates are discarded.
- = The test is performed but no candidate is discarded.

	Fold	Alive	Best	Mean best	Exp so far
x	1	9	4	0.9543	9
=	2	9	3	0.9433	18
-	3	3	4	0.9567	27
-	4	2	4	0.9566	30
=	5	2	4	0.9582	32
=	6	2	4	0.9546	34
=	7	2	4	0.9531	36
=	8	2	4	0.9539	38
=	9	2	4	0.9531	40
=	10	2	4	0.9529	42

Selected candidate: **ubSMOTE**      metric: auc      mean value: 0.9529



# Racing for Strategy Selection

```
> results
$best
[1] "ubSMOTE"

$avg
[1] 0.9529177

$sd
[1] 0.009049014

$N.test
[1] 42

$Gain
[1] 53
```

```
$Race
```

	unbal	ubOver	ubUnder	ubSMOTE	ubOSS	ubCNN	ubENN	ubNCL	ubTomek
[1,]	0.8844582	0.9138946	0.9354739	0.9543104	0.8957273	0.9139340	0.9024656	0.9014143	0.9048642
[2,]	0.9116642	0.9104928	0.9511485	0.9507221	0.9037491	0.9104840	0.9139047	0.9094542	0.9105558
[3,]	0.8979478	0.9013642	0.9502417	0.9649361	0.9092505	0.9081796	0.9103668	0.9036617	0.9058917
[4,]	NA	NA	0.9503782	0.9564226	NA	NA	0.8999928	NA	NA
[5,]	NA	NA	0.9537802	0.9647722	NA	NA	NA	NA	NA
[6,]	NA	NA	0.9494913	0.9362763	NA	NA	NA	NA	NA
[7,]	NA	NA	0.9411979	0.9440379	NA	NA	NA	NA	NA
[8,]	NA	NA	0.9576971	0.9594249	NA	NA	NA	NA	NA
[9,]	NA	NA	0.9530119	0.9473722	NA	NA	NA	NA	NA
[10,]	NA	NA	0.9633438	0.9509024	NA	NA	NA	NA	NA

```
> # Race using 4 cores and 500 trees (default)
> results <- ubRacing(Class ~., creditcard, "randomForest",
                      positive=1, metric="auc", ubConf=ubConf, ncore=4)

> library(e1071)
> results <- ubRacing(Class ~., creditcard, "svm",
                      positive=1, ubConf=ubConf)

> library(rpart)
> results <- ubRacing(Class ~., creditcard, "rpart",
                      positive=1, ubConf=ubConf)
```

# Useful R Packages

70/70

**imbalance**: Preprocessing Algorithms for Imbalanced Datasets

<https://cran.r-project.org/web/packages/imbalance/index.html>

Working with imbalanced datasets

<https://cran.r-project.org/web/packages/imbalance/vignettes/imbalance.pdf>

## mlr: Machine Learning in R

<https://cran.r-project.org/web/packages/mlr/vignettes/mlr.html>

### mlr: Machine Learning in R


Bernd Bischl, Michel Lang, Jakob Richter, Jakob Bossek, Leonard Judd, Tobias Kuehn, Erich Studerus, Lars Kotthoff

2017-03-14

### mlr: Machine Learning in R


This Vignette is supposed to give you a short introductory glance at the key features of `mlr`. A more detailed in depth and continuously updated tutorial can be found on the GitHub project page:

- [Project Page](#)
- Online Tutorial for [mlr release](#) and [mlr devel](#)
- Download the online tutorial for [mlr release](#) and [mlr devel](#) as zip for offline usage



## DMwR: Functions and data for "Data Mining with R"

<https://cran.r-project.org/web/packages/DMwR/index.html>



## XGBoost: eXtreme Gradient Boosting

(used for supervised learning tasks such as Regression, Classification, and Ranking)

<https://github.com/dmlc/xgboost>

<http://xgboost.readthedocs.io/en/latest/>

How to use XGBoost algorithm in R in easy steps

<https://www.analyticsvidhya.com/blog/2016/01/xgboost-algorithm-easy-steps/>

Kaggle 神器 XGBoost 入門：為什麼要用它?怎麼用?

<https://weiwenku.net/d/100778240>

# 進階選讀

# Simple Moving Average

- A moving average (移動平均) (簡稱均線) is a calculation to analyze data points by creating series of **averages** of **different subsets** of the full data set.

$$SMA = \frac{p_1 + p_2 + \cdots + p_n}{n}$$

**When price is in an uptrend** and subsequently, the moving average is in an uptrend, and the moving average has been tested by price and price has bounced off the moving average a few times (i.e. the moving average is serving as a support line), then a trader might **buy** on the **next pullbacks back** to the Simple Moving Average.

## Moving Average Acting as Support - Potential Buy Signal



www.OnlineTradingConcepts.com - All Rights Reserved

Created with TradeStation

<http://www.onlinetradingconcepts.com/TechnicalAnalysis/MASimple.html>

# Moving Average Acting as Resistance - Potential Sell Signal

73/70

- At times when price is in a **downtrend** and the moving average is in a downtrend as well, and price tests the SMA above and is rejected a few consecutive times (i.e. the moving average is serving as a resistance line), then a trader might sell on the **next rally up** to the Simple Moving Average.



## An n-day WMA (Weighted moving average)

$$WMA_M = \frac{np_M + (n-1)p_{M-1} + \cdots + 2p_{(M-n+2)} + p_{(M-n+1)}}{n + (n-1) + \cdots + 2 + 1}$$

www.OnlineTradingConcepts.com - All Rights Reserved

Created with TradeStation

<http://www.onlinetradingconcepts.com/TechnicalAnalysis/MASimple.html>

# Smoothing in R

## **smooth:** Forecasting Using Smoothing Functions

<https://cran.r-project.org/web/packages/smooth/index.html>

**es()** - Exponential Smoothing;  
**ssarima()** - State-Space ARIMA, also known as Several Seasonalities ARIMA;  
**ces()** - Complex Exponential Smoothing;  
**ges()** - Generalised Exponential Smoothing;  
**ves()** - Vector Exponential Smoothing;  
**sma()** - Simple Moving Average in state-space form;

## **TTR:** Technical Trading Rules

<https://cran.r-project.org/web/packages/TTR/index.html>

**SMA(x, n = 10, ...)**  
**EMA(x, n = 10, wilder = FALSE, ratio = NULL, ...)**  
**DEMA(x, n = 10, v = 1, wilder = FALSE, ratio = NULL)**  
**WMA(x, n = 10, wts = 1:n, ...)**  
**EVWMA(price, volume, n = 10, ...)**  
**ZLEMA(x, n = 10, ratio = NULL, ...)**  
**VWAP(price, volume, n = 10, ...)**  
**VMA(x, w, ratio = 1, ...)**  
**HMA(x, n = 20, ...)**  
**ALMA(x, n = 9, offset = 0.85, sigma = 6, ...)**

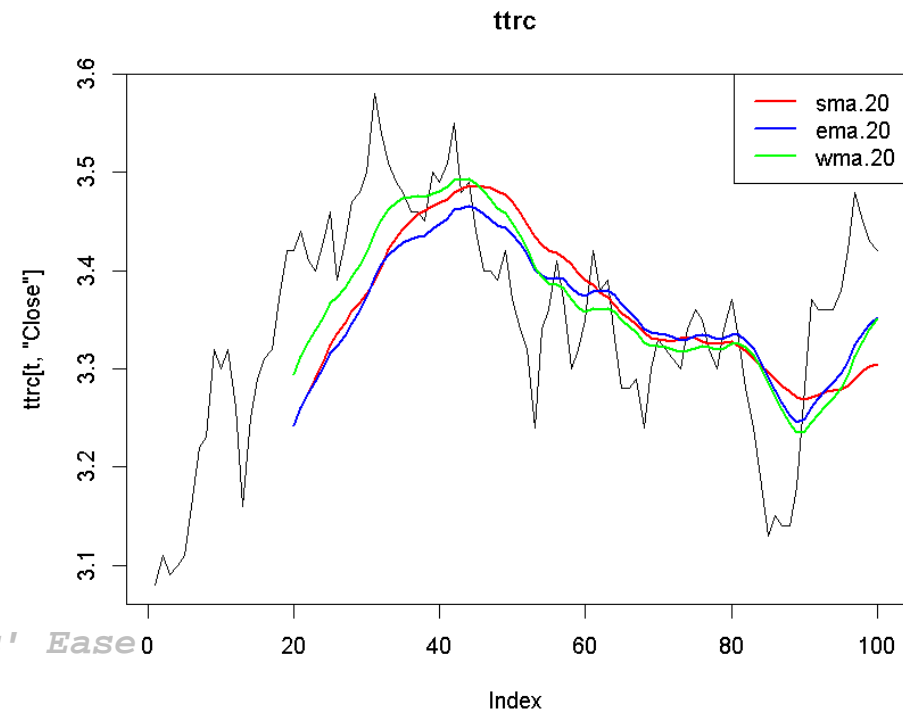
# Example

**ttrc {TTR}**: Technical Trading Rule Composite data

Historical Open, High, Low, Close, and Volume data for the periods January 2, 1985 to December 31, 2006. Randomly generated.

```
> # install.packages("TTR")
> library(TTR)
> data(ttrc)
> dim(ttrc)
[1] 5550    6
> head(ttrc)
      Date Open High  Low Close Volume
1 1985-01-02 3.18 3.18 3.08  3.08 1870906
2 1985-01-03 3.09 3.15 3.09  3.11 3099506
3 1985-01-04 3.11 3.12 3.08  3.09 2274157
4 1985-01-07 3.09 3.12 3.07  3.10 2086758
5 1985-01-08 3.10 3.12 3.08  3.11 2166348
6 1985-01-09 3.12 3.17 3.10  3.16 3441798

> t <- 1:100
> sma.20 <- SMA(ttrc[t, "Close"], 20)
> ema.20 <- EMA(ttrc[t, "Close"], 20) # Arms' Ease
> wma.20 <- WMA(ttrc[t, "Close"], 20)
>
> plot(ttrc[t, "Close"], type="l", main="ttrc")
> lines(sma.20, col="red", lwd=2)
> lines(ema.20, col="blue", lwd=2)
> lines(wma.20, col="green", lwd=2)
> legend("topright", legend=c("sma.20", "ema.20", "wma.20"),
+       col=c("red", "blue", "green"), lty=1, lwd=2)
```






# 曲線配適 (Fitting Curves)

76/70


## Example Methods

- non-linear parametric curves
- lowess (a non-parametric curve fitter)
- loess (a modelling tool)
- gam (fits generalized additive models)
- lm (linear model)



## locally-weighted polynomial regression

Loess regression  
(locally weighted polynomial regression)






```
> data(cars)
> dim(cars)
[1] 50 2
```

```
> head(cars)
  speed dist
```

```
1      4     2
2      4    10
3      7     4
4      7    22
5      8    16
6      9    10
```

```
> par(mfrow=c(1, 3))
> for(i in c(0.1, 0.3, 0.5)){
+   plot(cars$dist ~ cars$speed, main=paste0("lowess (f=", i, ")"))
+   lines(lowess(cars$dist ~ cars$speed, f = i), col="red", lwd=2)
+ }
```




`cars {datasets}`:

The data give the speed of cars and the distances taken to stop.  
Note that the data were recorded in the 1920s.


# Density Plots (Smoothed Histograms) (1/3)

## Constructing a Smoothed Histogram (Jacoby, 1997)


A. Unidimensional scatterplot of 10 data points




B. Data points shown as kernel densities




histogram



C. Summing kernel densities at the 20  $v_i$



D. Final smoothed histogram

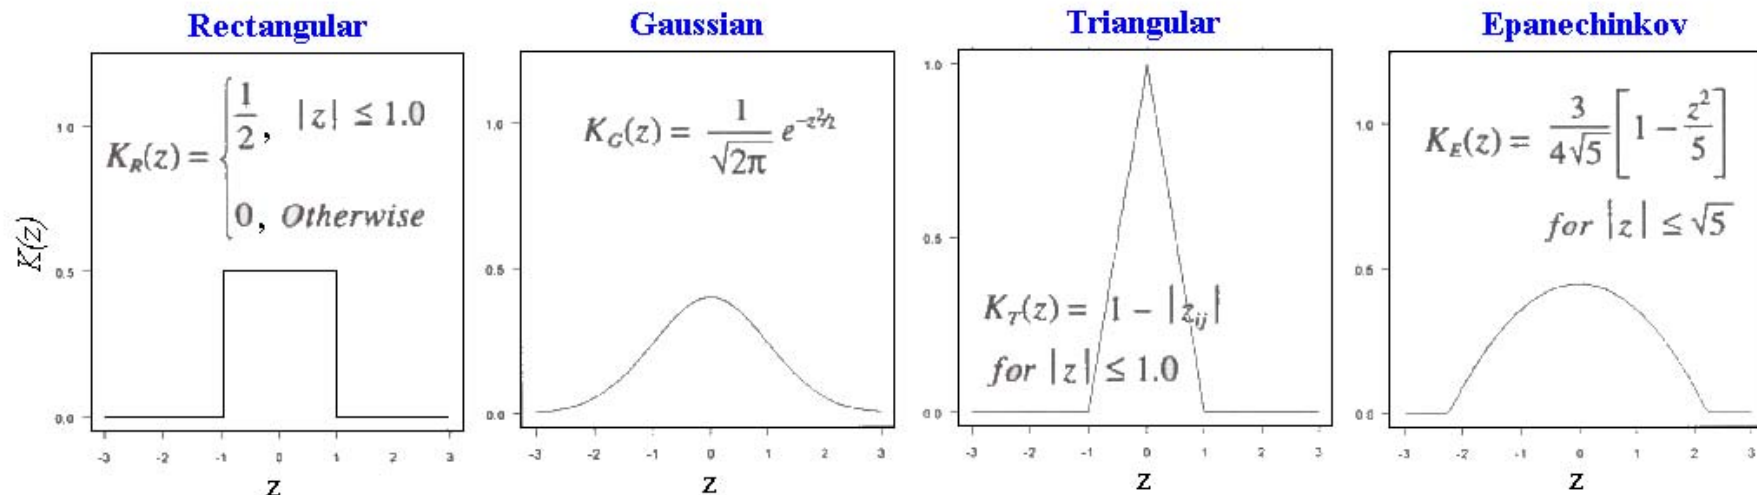


# Kernel Density Estimation

79/70

- Selection of **kernels ( $K$ )**
- Selection of **bandwidth ( $h$ )**

Figures modified from Jacoby (1997)



nonparametric regression

$$y_i = f_0(x_i) + \epsilon_i, \quad i = 1, \dots, n,$$

$\epsilon_1, \dots, \epsilon_n$  are still i.i.d. random errors with  $\mathbb{E}(\epsilon_i) = 0$

$$\hat{f}(v_i) = \frac{1}{hn} \sum_{i=1}^n K[z_{ij}]$$

$$z_{ij} = \frac{1}{h}(v_j - x_i)$$

kernel regression

$$\hat{f}(x) = \frac{\sum_{i=1}^n K\left(\frac{x - x_i}{h}\right) y_i}{\sum_{i=1}^n K\left(\frac{x - x_i}{h}\right)}$$


# Kernel Density Estimation

80/70


進階選讀

## Different kernels


Rectangular




Gaussian







Triangular



Epanechnikov



## Different bandwidth



# Kernel Density Estimation in R

81/70








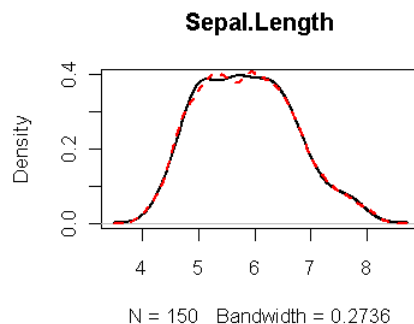
進階選讀

— gaussian  
- - - epanechnikov

```
density(x, bw = "nrd0", adjust = 1,  
        kernel = c("gaussian", "epanechnikov", "rectangular",  
                   "triangular", "biweight",  
                   "cosine", "optcosine"),  
        weights = NULL, window = kernel, width,  
        give.Rkern = FALSE,  
        n = 512, from, to, cut = 3, na.rm = FALSE, ...)
```

```
> plot(density(iris$Sepal.Length))
```

```
> hist(iris$Sepal.Length, prob=T)  
> lines(density(iris$Sepal.Length), col="red")
```






# Spline approximate to the top profile of the ruddy duck

82/70

進階選讀

## Ruddy duck (棕硬尾鴨) (雄)

"棕硬尾鴨棲息在北美洲的沼澤湖及池中，在南美洲的安地斯山脈也有分布。" <https://zh.wikipedia.org/wiki/棕硬尾鴨>



$x$	0.9	1.3	1.9	2.1	2.6	3.0	3.9	4.4	4.7	5.0	6.0	7.0	8.0	9.2	10.5	11.3	11.6	12.0	12.6	13.0	13.3
$f(x)$	1.3	1.5	1.85	2.1	2.6	2.7	2.4	2.15	2.05	2.1	2.25	2.3	2.25	1.95	1.4	0.9	0.7	0.6	0.5	0.4	0.25




## 三次樣條插值法

### Usage

```
smooth.spline(x, y = NULL, w = NULL, df, spar = NULL, lambda = NULL, cv = FALSE,  
              all.knots = FALSE, nknots = .nknots.smspl,  
              keep.data = TRUE, df.offset = 0, penalty = 1,  
              control.spar = list(), tol = 1e-6 * IQR(x), keep.stuff = FALSE)
```

```
> #install.packages("jpeg")  
> library(jpeg)  
> ruddyduck.img <- readJPEG("ruddyduck.jpg")  
> plot(0, xlim=c(0, 14), ylim=c(-6, 4), type='n', xlab="", ylab="",  
+      main="Spline approximate to the top profile of the ruddy duck")  
> rasterImage(ruddyduck.img, 0.6, -6, 13.8, 3.3)  
> abline(v=1:14, h=-6:4, col=gray)
```

Spline approximate to the top profile of the ruddy duck




# smooth.spline {stats}: Fit a Smoothing Spline

84/70


進階選讀

```
> ruddyduck.dat <- read.table("ruddyduck.txt", header=T, sep="\t")
> head(ruddyduck.dat)
      x    fx
1 0.9 1.30
2 1.3 1.50
3 1.9 1.85
4 2.1 2.10
5 2.6 2.60
6 3.0 2.70
> points(ruddyduck.dat, col="blue", pch=16)
>
> duck.spl <- smooth.spline(ruddyduck.dat$fx ~ ruddyduck.dat$x)
> lines(duck.spl, col = "red", lwd=2)
```

Spline approximate to the top profile of the ruddy duck



Spline approximate to the top profile of the ruddy duck



# Cubic Spline Interpolation

85/70

進階選讀

## Cubic Splines Interpolant

### Definition 3.10

Given a function  $f$  defined on  $[a, b]$  and a set of nodes  $a = x_0 < x_1 < \cdots < x_n = b$ , a cubic spline interpolant  $S$  for  $f$  is a function that satisfies the following conditions:

- (a)  $S(x)$  is a cubic polynomial ( $S_j(x)$ ) on  $[x_j, x_{j+1}]$ .
- (b)  $S_j(x_j) = f(x_j)$  and  $S_j(x_{j+1}) = f(x_{j+1})$ ,  $j = 0, 1, \dots, n-1$ ;
- (c)  $S_{j+1}(x_{j+1}) = S_j(x_{j+1})$ ; (d)  $S'_{j+1}(x_{j+1}) = S'_j(x_{j+1})$ ;
- (e)  $S''_{j+1}(x_{j+1}) = S''_j(x_{j+1})$  for each  $j = 0, 1, \dots, n-2$ ;
- (f) One of the following sets of boundary conditions is satisfied:
  - (i)  $S''(x_0) = S''(x_n) = 0$  (natural or free boundary);
  - (ii)  $S'(x_0) = f'(x_0)$  and  $S'(x_n) = f'(x_n)$  (clamped boundary).

## Construction of a Cubic Spline (conti.)

- (12) This system involves only the  $\{c_j\}_{j=0}^n$  as unknowns.
- (13) The values of  $\{h_j\}_{j=0}^{n-1}$  and  $\{a_j\}_{j=0}^n$  are given, respectively, by the spacing of the nodes  $\{x_j\}_{j=0}^n$  and the values of  $f$  at the nodes.
- (14) So once the values of  $\{c_j\}_{j=0}^n$  are determined, it is a simple matter to find the remainder of the constants  $\{b_j\}_{j=0}^{n-1}$  from Eq. (3.20) and  $\{d_j\}_{j=0}^{n-1}$  from Eq. (3.17).
- (15) The major question that arises in connection with this construction is whether the values of  $\{c_j\}_{j=0}^n$  can be found using the system of equations given in (3.21) and, if so, whether these values are unique.

## ALGORITHM 034: Natural Cubic Spline

To construct the cubic spline interpolant  $S$  for the function  $f$ , defined at the numbers  $x_0 < x_1 < \cdots < x_n$ , satisfying  $S''(x_0) = S''(x_n) = 0$ :

INPUT  $n; x_0, x_1, \dots, x_n; a_0 = f(x_0), a_1 = f(x_1), \dots, a_n = f(x_n)$ .

OUTPUT  $a_j, b_j, c_j, d_j$  for  $j = 0, 1, \dots, n-1$ .

(Note:  $S(x) = S_j(x) = a_j + b_j(x - x_j) + c_j(x - x_j)^2 + d_j(x - x_j)^3$  for  $x_j \leq x \leq x_{j+1}$ .)

Step 1 For  $i = 0, 1, \dots, n-1$  set  $h_i = x_{i+1} - x_i$ .

Step 2 For  $i = 1, 2, \dots, n-1$  set

$$\alpha_i = \frac{3}{h_i}(a_{i+1} - a_i) - \frac{3}{h_{i-1}}(a_i - a_{i-1}).$$

Step 3 Set  $l_0 = 1$ ; (Steps 3, 4, 5, and part of Step 6 solve a tridiagonal linear system using a method described in Algorithm 6.7.)

$$\mu_0 = 0;$$
$$z_0 = 0.$$

## ALGORITHM 034: Natural Cubic Spline (conti.)

Step 4 For  $i = 1, 2, \dots, n-1$

set  $l_i = 2(x_{i+1} - x_{i-1}) - h_{i-1}\mu_{i-1}$ ;

$$\mu_i = h_i/l_i;$$

$$z_i = (\alpha_i - h_{i-1}z_{i-1})/l_i.$$

Step 5 Set  $l_n = 1$ ;

$$z_n = 0;$$

$$c_n = 0.$$

Step 6 For  $j = n-1, n-2, \dots, 0$

set  $c_j = z_j - \mu_j c_{j+1}$ ;

$$b_j = (a_{j+1} - a_j)/h_j - h_j(c_{j+1} + 2c_j)/3;$$

$$d_j = (c_{j+1} - c_j)/(3h_j).$$

Step 7 OUTPUT  $(a_j, b_j, c_j, d_j)$  for  $j = 0, 1, \dots, n-1$ ;

STOP.